



Logistic Regression

Industrial AI Lab.


Linear Classification: Logistic Regression

- Logistic regression is a classification algorithm
 - don't be confused
- Perceptron: make use of sign of data
- SVM: make use of margin (minimum distance)
 - Distance from a single data point
- We want to use distance information of **all** data points
 - logistic regression


Using Distances



Using Distances




$$|h_1| + |h_2|$$




$$|h_1| + |h_2|$$

Using Distances



$$|h_1| + |h_2|$$



$$|h_1| \cdot |h_2|$$



$$|h_1| + |h_2|$$

$$|h_1| \cdot |h_2|$$

Using Distances



$$|h_1| + |h_2|$$

$$|h_1| + |h_2|$$


$$|h_1| \cdot |h_2|$$

$$|h_1| \cdot |h_2|$$

$$\frac{|h_1| + |h_2|}{2} \geq \sqrt{|h_1| \cdot |h_2|} \quad \text{equal iff } |h_1| = |h_2|$$

Using all Distances

- basic idea: to find the decision boundary (hyperplane) of $g(x) = \omega^T x = 0$ such that maximizes $\prod_i |h_i| \rightarrow$ **optimization**




- Inequality of arithmetic and geometric means

$$\frac{x_1 + x_2 + \cdots + x_m}{m} \geq \sqrt[m]{x_1 \cdot x_2 \cdots x_m}$$

and that equality holds if and only if $x_1 = x_2 = \cdots = x_m$

Using all Distances




- Roughly speaking, this optimization of $\max \prod_i |h_i|$ tends to position a **hyperplane in the middle of two classes**


$$h = \frac{g(x)}{\|\omega\|} = \frac{\omega^T x}{\|\omega\|} \sim \omega^T x$$

Using all Distances with Outliers

- SVM vs. Logistic Regression




SVM






Logistic Regression

Sigmoid Function

- We link or squeeze $(-\infty, +\infty)$ to $(0, 1)$ for several reasons:



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$




Sigmoid Function

```
import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
```

```
z = np.linspace(-10,10,100)
s = 1/(1+np.exp(-z))

plt.figure(figsize=(10,6))
plt.plot(z, s)
plt.xlim([-10, 10])
plt.ylim([-0.1, 1.1])
plt.show()
```



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Sigmoid Function

- $\sigma(z)$ is the sigmoid function, or the logistic function
 - Logistic function always generates a value between 0 and 1
 - Crosses 0.5 at the origin, then flattens out

$$\sigma(z) = \frac{1}{1 + e^{-z}} \implies \sigma(\omega^T x) = \frac{1}{1 + e^{-\omega^T x}}$$

Sigmoid Function

- Benefit of mapping via the logistic function
 - Monotonic: same or similar optimization solution
 - Continuous and differentiable: good for gradient descent optimization
 - Probability or confidence: can be considered as probability

$$P(y = +1 \mid x, \omega) = \frac{1}{1 + e^{-\omega^T x}} \in [0, 1]$$

- Often we do not care about predicting the label y
- Rather, we want to predict the label probabilities $P(y \mid x, \omega)$
 - Probability that the label is +1

$$P(y = +1 \mid x, \omega)$$

- Probability that the label is 0

$$P(y = 0 \mid x, \omega) = 1 - P(y = +1 \mid x, \omega)$$

- Goal: we need to fit ω to our data

Probabilistic Approach (or MLE)

- Consider a random variable $y \in \{0, 1\}$

$$P(y = +1) = p, \quad P(y = 0) = 1 - p$$

where $p \in [0, 1]$, and is assumed to depend on a vector of explanatory variables $x \in \mathbb{R}^n$

- Then, the logistic model has the form

$$p = \frac{1}{1 + e^{-\omega^T x}} = \frac{e^{\omega^T x}}{e^{\omega^T x} + 1}$$
$$1 - p = \frac{1}{e^{\omega^T x} + 1}$$

- We can re-order the training data so
 - for x_1, \dots, x_q , the outcome is $y = +1$, and
 - for x_{q+1}, \dots, x_m , the outcome is $y = 0$

Probabilistic Approach (or MLE)

- Likelihood function

$$\mathcal{L} = \prod_{i=1}^q p_i \prod_{i=q+1}^m (1 - p_i) \quad \left(\sim \prod_i |h_i| \right)$$

- Log likelihood function

$$\begin{aligned} \ell(\omega) &= \log \mathcal{L} = \sum_{i=1}^q \log p_i + \sum_{i=q+1}^m \log(1 - p_i) \\ &= \sum_{i=1}^q \log \frac{\exp(\omega^T x_i)}{1 + \exp(\omega^T x_i)} + \sum_{i=q+1}^m \log \frac{1}{1 + \exp(\omega^T x_i)} \\ &= \sum_{i=1}^q (\omega^T x_i) - \sum_{i=1}^m \log(1 + \exp(\omega^T x_i)) \end{aligned}$$

- Since ℓ is a concave function of ω , the logistic regression problem can be solved as a convex optimization problem

$$\hat{\omega} = \arg \max_{\omega} \ell(\omega)$$

In Matrix Form

$$\omega = \begin{bmatrix} \omega_1 \\ \omega_2 \\ \omega_3 \end{bmatrix}, \quad x = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix}$$

$$X = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ (x^{(3)})^T \\ \vdots \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} \\ 1 & x_1^{(3)} & x_2^{(3)} \\ \vdots & \vdots & \vdots \end{bmatrix}$$

Source: Section 7.1.1

from http://cvxr.com/cvx/examples/cvxbook/Ch07_statistical_estim/html/logistics.html 16

Data Generation

```
m = 100
w = np.array([[ -4], [2], [1]])
X = np.hstack([np.ones([m,1]), 2*np.random.rand(m,1), 4*np.random.rand(m,1)])


w = np.asmatrix(w)
X = np.asmatrix(X)

y = (np.exp(X*w)/(1+np.exp(X*w))) > 0.5

C1 = np.where(y == True)[0]
C2 = np.where(y == False)[0]

y = np.empty([m,1])
y[C1] = 1
y[C2] = 0
y = np.asmatrix(y)

plt.figure(figsize = (10,6))
plt.plot(X[C1,1], X[C1,2], 'ro', label='C1')
plt.plot(X[C2,1], X[C2,2], 'bo', label='C2')
plt.legend()
plt.show()
```



Log Likelihood

$$\begin{aligned}\ell(\omega) &= \log \mathcal{L} = \sum_{i=1}^q \log p_i + \sum_{i=q+1}^m \log(1 - p_i) \\ &= \sum_{i=1}^q \log \frac{\exp(\omega^T x_i)}{1 + \exp(\omega^T x_i)} + \sum_{i=q+1}^m \log \frac{1}{1 + \exp(\omega^T x_i)} \\ &= \sum_{i=1}^q (\omega^T x_i) - \sum_{i=1}^m \log(1 + \exp(\omega^T x_i))\end{aligned}$$

- Refer to [cvx functions](#)

- scalar function: `cvx.sum_entries(x)` = $\sum_{ij} x_{ij}$
- elementwise function: `cvx.logistic(x)` = $\log(1 + e^x)$

CVXPY

```
import cvxpy as cvx
```

```
w = cvx.Variable(3, 1)
```

```
obj = cvx.Maximize(y.T*X*w - cvx.sum_entries(cvx.logistic(X*w)))  
prob = cvx.Problem(obj).solve()
```


```
w = w.value
```

```
xp = np.linspace(0,2,100).reshape(-1,1)  
yp = - w[1,0]/w[2,0]*xp - w[0,0]/w[2,0]
```

```
plt.figure(figsize = (10,6))  
plt.plot(X[C1,1], X[C1,2], 'ro', label='C1')  
plt.plot(X[C2,1], X[C2,2], 'bo', label='C2')  
plt.plot(xp, yp, 'k', label='Logistic Regression')  
plt.legend()  
plt.show()
```

$$\text{cvx.logistic}(x) = \log(1 + e^x)$$

$$\sum_{i=1}^q (\omega^T x_i) - \sum_{i=1}^m \log(1 + \exp(\omega^T x_i))$$



In a More Compact Form

- Change $y \in \{0, +1\} \rightarrow y \in \{-1, +1\}$ for computational convenience

– Consider the following function

$$P(y = +1) = p = \sigma(\omega^T x), \quad P(y = -1) = 1 - p = 1 - \sigma(\omega^T x) = \sigma(-\omega^T x)$$
$$P(y \mid x, \omega) = \sigma(y\omega^T x) = \frac{1}{1 + \exp(-y\omega^T x)} \in [0, 1]$$

– Log-likelihood

$$\begin{aligned} \ell(\omega) = \log \mathcal{L} = \log P(y \mid x, \omega) &= \log \prod_{n=1}^m P(y_n \mid x_n, \omega) \\ &= \sum_{n=1}^m \log P(y_n \mid x_n, \omega) \\ &= \sum_{n=1}^m \log \frac{1}{1 + \exp(-y_n \omega^T x_n)} \\ &= \sum_{n=1}^m -\log(1 + \exp(-y_n \omega^T x_n)) \end{aligned}$$

In a More Compact Form

- MLE solution

$$\begin{aligned}\hat{\omega} &= \arg \max_{\omega} \sum_{n=1}^m -\log(1 + \exp(-y_n \omega^T x_n)) \\ &= \arg \min_{\omega} \sum_{n=1}^m \log(1 + \exp(-y_n \omega^T x_n))\end{aligned}$$

CVXPY

```
y = np.empty([m,1])
y[C1] = 1
y[C2] = -1
y = np.asmatrix(y)
```


```
w = cvx.Variable(3, 1)
```

```
obj = cvx.Minimize(cvx.sum_entries(cvx.logistic(-cvx.mul_elemwise(y,X*w))))
prob = cvx.Problem(obj).solve()
```

```
w = w.value
```


$$\text{cvx.logistic}(x) = \log(1 + e^x)$$

$$= \arg \min_{\omega} \sum_{n=1}^m \log(1 + \exp(-y_n \omega^T x_n))$$



Logistic Regression

- Classified based on probability



Multiclass Classification

- Generalization to more than 2 classes is straightforward
 - one vs. all (one vs. rest)
 - one vs. one
- Using the soft-max function instead of the logistic function
 - (refer to [UFLDL Tutorial](#))
 - see them as probability


$$P(y = k \mid x, \omega) = \frac{\exp(\omega_k^T x)}{\sum_k \exp(\omega_k^T x)} \in [0, 1]$$

- We maintain a separator weight vector ω_k for each class k

Non-linear Classification

- Same idea as non-linear regression: non-linear features
 - Explicit or implicit Kernel

```
X1 = np.array([[ -1.1, 0], [ -0.3, 0.1], [ -0.9, 1], [ 0.8, 0.4], [ 0.4, 0.9], [ 0.3, -0.6],  
              [-0.5, 0.3], [-0.8, 0.6], [-0.5, -0.5]])  
  
X2 = np.array([[ -1, -1.3], [ -1.6, 2.2], [ 0.9, -0.7], [ 1.6, 0.5], [ 1.8, -1.1], [ 1.6, 1.6],  
              [-1.6, -1.7], [-1.4, 1.8], [ 1.6, -0.9], [ 0, -1.6], [ 0.3, 1.7], [-1.6, 0], [-2.1, 0.2]])  
  
X1 = np.asmatrix(X1)  
X2 = np.asmatrix(X2)  
  
plt.figure(figsize=(10, 6))  
plt.plot(X1[:, 0], X1[:, 1], 'ro', label='C1')  
plt.plot(X2[:, 0], X2[:, 1], 'bo', label='C2')  
plt.axis([-3, 3, -3, 3])  
plt.legend(loc = 4, fontsize = 15)  
plt.show()
```



Explicit Kernel

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \implies z = \phi(x) = \begin{bmatrix} 1 \\ \sqrt{2}x_1 \\ \sqrt{2}x_2 \\ x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

```
N = X1.shape[0]
M = X2.shape[0]

X = np.vstack([X1, X2])
y = np.vstack([np.ones([N,1]), -np.ones([M,1])])

X = np.asmatrix(X)
y = np.asmatrix(y)

m = N + M
Z = np.hstack([np.ones([m,1]), np.sqrt(2)*X[:,0], np.sqrt(2)*X[:,1], np.square(X[:,0]),
               np.sqrt(2)*np.multiply(X[:,0],X[:,1]), np.square(X[:,1])])

w = cvx.Variable(6, 1)
obj = cvx.Minimize(cvx.sum_entries(cvx.logistic(-cvx.mul_elemwise(y,Z*w))))
prob = cvx.Problem(obj).solve()

w = w.value
```

Non-linear Classification

