

# Parameter Estimation in Probabilistic Model

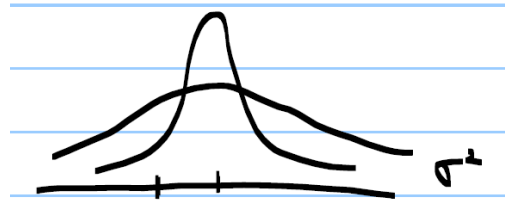
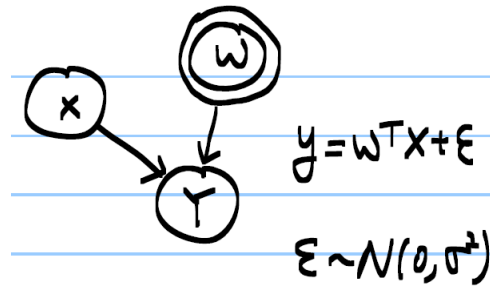
by Prof. Seungchul Lee  
iSystems Design Lab  
<http://isystems.unist.ac.kr/>  
UNIST

## Table of Contents

- I. 1. Generative model
- II. 2. Maximum Likelihood Estimation (MLE)
  - I. 2.1. Given  $m$  data points, drawn from an exponential distribution
  - II. 2.2. Given  $m$  data points, drawn from a Gaussian distribution
  - III. 2.3. Linear Regression: A Probabilistic View
  - IV. 2.4. Data Fusion with Uncertainties
- III. 3. Maximum-a-Posterior Estimation (MAP)
  - I. 3.1. MAP for mean of a univariate Gaussian,  $\mathcal{N}(\theta, \sigma^2)$
  - II. 3.2. MAP Python code
  - III. 3.3. Object Tracking in Computer Vision
- IV. 4. Kernel Density Estimation

# 1. Generative model

$$P(y | X, \omega, \sigma^2) = \mathcal{N}(\omega^T X, \sigma^2)$$



## 2. Maximum Likelihood Estimation (MLE)

Estimate parameters  $\theta (\omega, \sigma^2)$  such that maximize the likelihood given a generative model

- Given observed data

$$D = \{(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)\}$$

- Generative model structure

$$\begin{aligned} y_i &= \hat{y}_i + \varepsilon \\ &= \omega^T x_i + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2) \end{aligned}$$

- Find parameters  $\omega$  and  $\sigma$  that maximize the likelihood over the observed data
- Likelihood:

$$\begin{aligned} \mathcal{L}(\omega, \sigma) &= P \left( y_1, y_2, \dots, y_m \mid x_1, x_2, \dots, x_m; \underbrace{\omega, \sigma}_{\theta} \right) \\ &= \prod_{i=1}^m P(y_i \mid x_i; \omega, \sigma) \\ &= \frac{1}{(2\pi\sigma^2)^{\frac{m}{2}}} \exp \left( -\frac{1}{2\sigma^2} \sum_{i=1}^m (y_i - \omega^T x_i)^2 \right) \end{aligned}$$

- Perhaps the simplest (but widely used) parameter estimation method

### 2.1. Given $m$ data points, drawn from an exponential distribution

- Exponential distribution from [fundamentals of statistics](http://www.statlect.com/exponential_distribution_maximum_likelihood.htm)  
([http://www.statlect.com/exponential\\_distribution\\_maximum\\_likelihood.htm](http://www.statlect.com/exponential_distribution_maximum_likelihood.htm))

$$f(y) = \frac{1}{a} \exp \left( -\frac{1}{a} y \right) : \text{generative model}$$

$$\begin{aligned} \mathcal{L} &= P(y_1, y_2, \dots, y_m \mid a) \\ &= \prod_{i=1}^m \frac{1}{a} \exp \left( -\frac{1}{a} y_i \right) \\ &= \frac{1}{a^m} \exp \left( -\frac{1}{a} \sum_{i=1}^m y_i \right) \end{aligned}$$

$$\text{Log-likelihood } \ell = \log \mathcal{L} = -m \log a - \frac{1}{a} \sum_{i=1}^m y_i$$

- Find  $a$  that maximizes  $\ell$

$$\frac{d\ell}{da} = -\frac{m}{a} + \frac{1}{a^2} \sum_{i=1}^m y_i = 0$$

$$\therefore a_{ML} = \frac{1}{m} \sum_{i=1}^m y_i \quad : \text{sample mean}$$

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import norm

%matplotlib inline
```

In [2]:

```
# exponential random variable
m = 50
x = np.random.exponential(700, (m, 1)) # mu = 700

# MLE
print(1/m*np.sum(x))
```

514.030918206

## 2.2. Given $m$ data points, drawn from a Gaussian distribution

$$P(y = y_i \mid \mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(y_i - \mu)^2\right) : \text{generative model}$$

$$\begin{aligned} \mathcal{L} = P(y_1, y_2, \dots, y_m \mid \mu, \sigma^2) &= \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(y_i - \mu)^2\right) \\ &= \frac{1}{(2\pi)^{\frac{m}{2}} \sigma^m} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^m (y_i - \mu)^2\right) \end{aligned}$$

$$\ell = \log \mathcal{L} = -\frac{m}{2} \log 2\pi - m \log \sigma - \frac{1}{2\sigma^2} \sum_{i=1}^m (y_i - \mu)^2$$

- To maximize,  $\frac{\partial \ell}{\partial \mu} = 0, \frac{\partial \ell}{\partial \sigma} = 0$

$$\frac{\partial \ell}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^m (y_i - \mu) = 0 \implies \mu_{ML} = \frac{1}{m} \sum_{i=1}^m y_i \quad : \text{sample mean}$$

$$\frac{\partial \ell}{\partial \sigma} = -\frac{m}{\sigma} + \frac{1}{\sigma^3} \sum_{i=1}^m (y_i - \mu)^2 = 0 \implies \sigma_{ML}^2 = \frac{1}{m} \sum_{i=1}^m (y_i - \mu)^2 \quad : \text{sample variance}$$

- BIG Lesson
  - We often compute a mean and variance to represent data statistics
  - We kind of assume that a data set is Gaussian distributed
  - Good news: sample mean is Gaussian distributed by the central limit theorem

## Numerical Simulation

- Compute the likelihood function, then
  - maximize the likelihood function
  - adjust the mean and variance of the Gaussian to maximize its product

In [3]:

```
# MLE of Gaussian distribution
# mu

m = 20
mu = 0
sigma = 5

x = np.random.normal(mu, sigma, [m, 1])
xp = np.linspace(-20, 20, 100)
y0 = np.zeros([m, 1])

muhat = [-5, 0, 5, np.mean(x)]

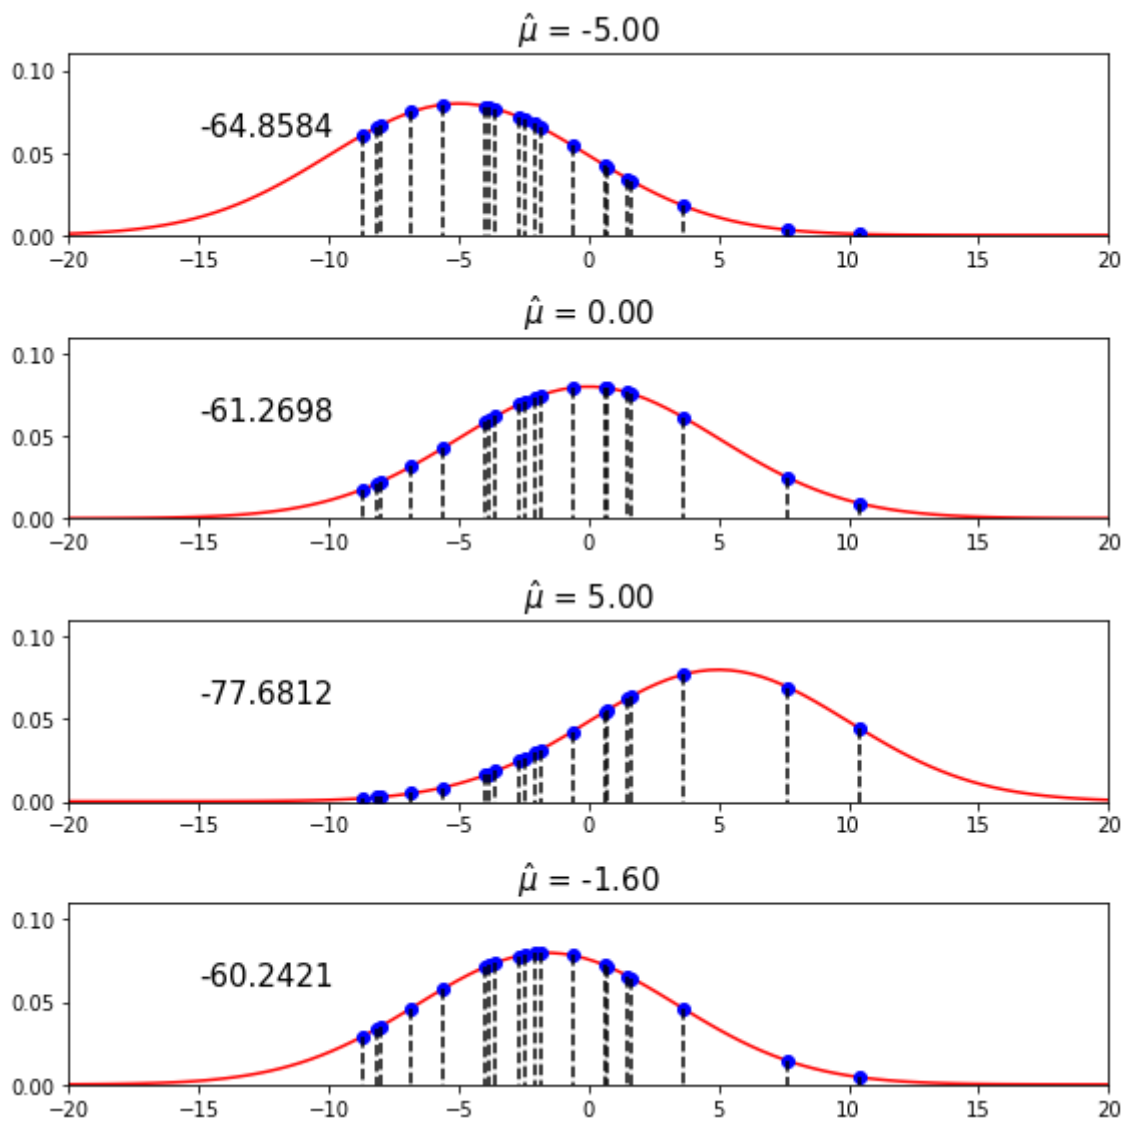
plt.figure(figsize=(8, 8))

for i in range(4):
    yp = norm.pdf(xp, muhat[i], sigma)
    y = norm.pdf(x, muhat[i], sigma)
    logL = np.sum(np.log(y))

    plt.subplot(4, 1, i+1)
    plt.plot(xp, yp, 'r')
    plt.plot(x, y, 'bo')
    plt.plot(np.hstack([x, x]).T, np.hstack([y, y0]).T, 'k--')

    plt.title(r'$\hat{\mu}$ = {0:.2f}'.format(muhat[i]), fontsize=15)
    plt.text(-15, 0.06, np.round(logL, 4), fontsize=15)
    plt.axis([-20, 20, 0, 0.11])

plt.tight_layout()
plt.show()
```



Compare to a result from formula

$$\mu_{ML} = \frac{1}{m} \sum_{i=1}^m x_i$$

In [5]:

```
# mean is unknown in this example
# variance is known in this example

m = 10
mu = 0
sigma = 5

x = np.random.normal(mu, sigma, [m,1])

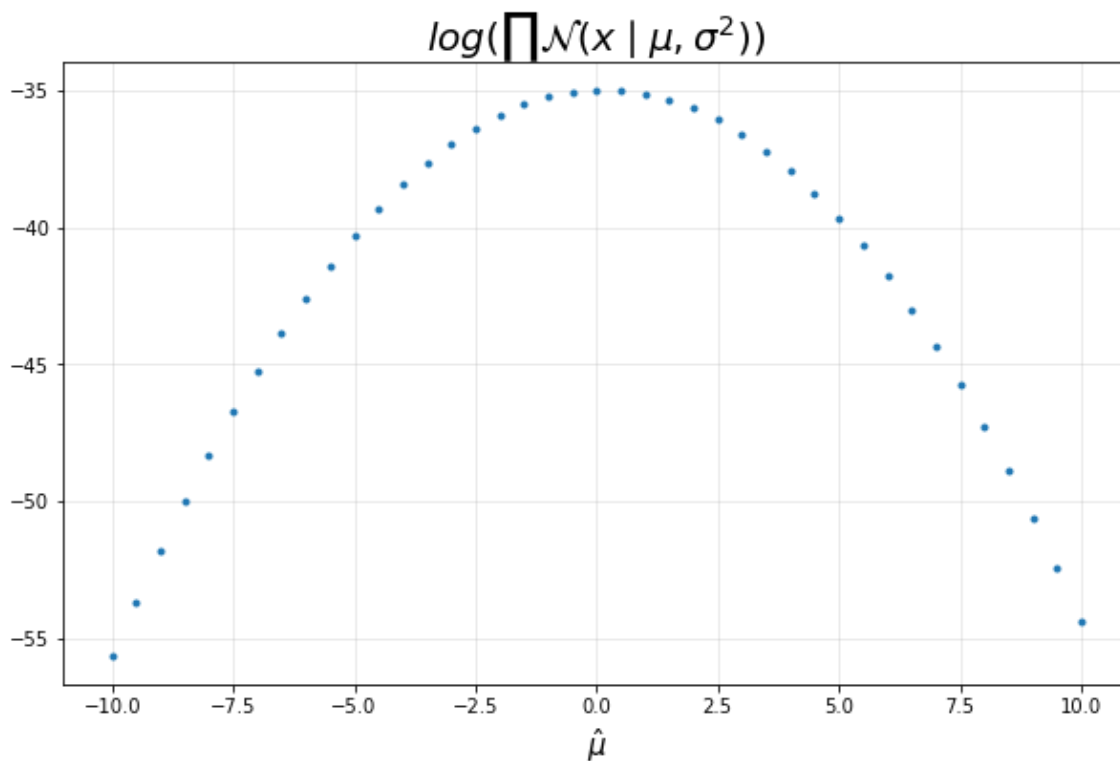
mus = np.arange(-10, 10.5, 0.5)
LOGL = []

for i in range(np.size(mus)):
    y = norm.pdf(x, mus[i], sigma)
    logL = np.sum(np.log(y))
    LOGL.append(logL)

muhat = np.mean(x)
print(muhat)

plt.figure(figsize=(10, 6))
plt.plot(mus, LOGL, '.')
plt.title('$\log(\prod \mathcal{N}(x \mid \mu, \sigma^2))$', fontsize=20)
plt.xlabel(r'$\hat{\mu}$', fontsize=15)
plt.grid(alpha=0.3)
plt.show()
```

0.160329485196



$$\sigma_{ML}^2 = \frac{1}{m} \sum_{i=1}^m (y_i - \mu)^2$$



In [6]:

```
# mean is known in this example
# variance is unknown in this example

m = 100
mu = 0
sigma = 3

x = np.random.normal(mu,sigma,[m,1]) # samples

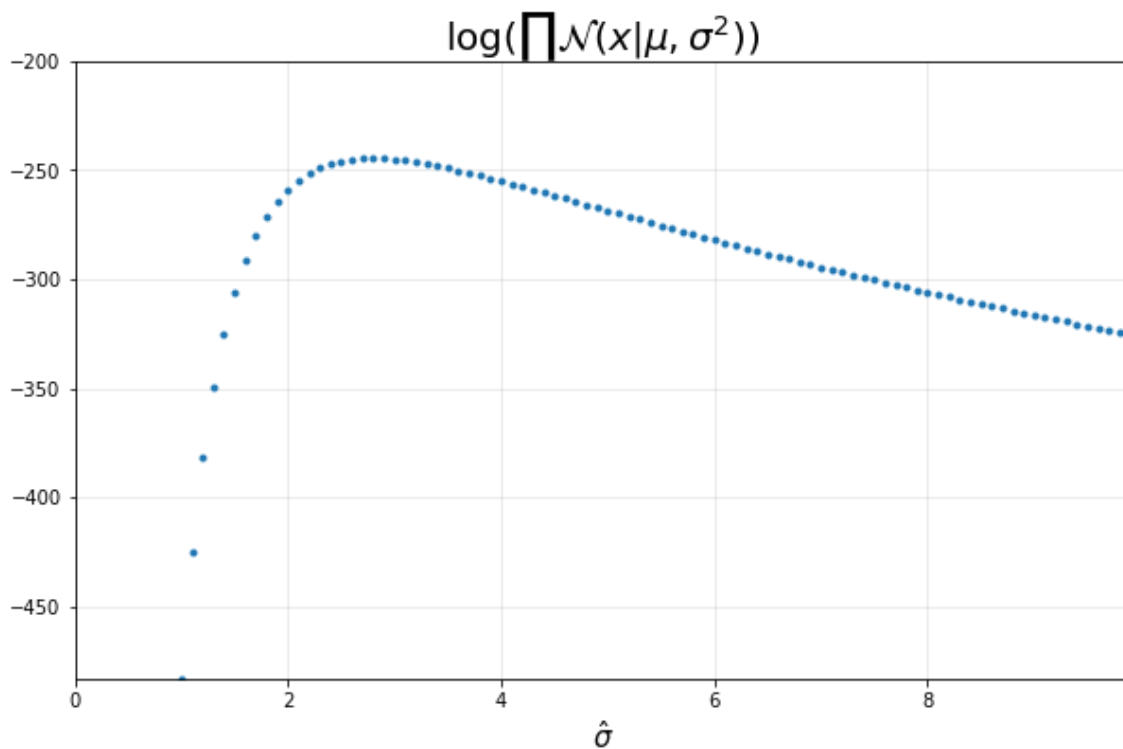
sigmas = np.arange(1, 10, 0.1)
LOGL = []

for i in range(sigmas.shape[0]):
    y = norm.pdf(x, mu, sigmas[i])    # likelihood
    logL = np.sum(np.log(y))
    LOGL.append(logL)

sigmahat = np.sqrt(np.var(x))
print(sigmahat)

plt.figure(figsize=(10,6))
plt.title(r'$\log (\prod \mathcal{N}(x|\mu,\sigma^2))$', fontsize=20)
plt.plot(sigmas, LOGL, '.')
plt.xlabel(r'$\hat{\sigma}$', fontsize=15)
plt.axis([0, np.max(sigmas), np.min(LOGL), -200])
plt.grid(alpha=0.3)
plt.show()
```

2.79684136967



## 2.3. Linear Regression: A Probabilistic View

- Probabilistic Machine Learning
  - I personally believe this is a more fundamental way of looking at machine learning
- Linear regression model with (Gaussian) normal errors

$$y = \omega^T x + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$
$$y - \omega^T x = \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

$$P(y_i | x_i; \omega, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2} (y_i - \omega^T x_i)^2\right): \text{generative model}$$

$$\begin{aligned}\mathcal{L} &= P(y_1, y_2, \dots, y_m | \omega, \sigma^2) \\ &= \prod_{i=1}^m P(y_i | x_i; \omega, \sigma^2) \\ &= \frac{1}{(\sqrt{2\pi})^m} \frac{1}{\sigma^m} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^m (y_i - \omega^T x_i)^2\right) = \text{likelihood}\end{aligned}$$

$$\ell = -\frac{m}{2} \log 2\pi - m \log \sigma - \frac{1}{2\sigma^2} \sum_{i=1}^m (y_i - \omega^T x_i)^2$$

$$\begin{aligned}\frac{d\ell}{d\omega} &= -2X^T Y + 2X^T X\omega = 0 \implies \omega_{ML} = (X^T X)^{-1} X^T Y \quad (\text{look familiar ?}) \\ \frac{d\ell}{d\sigma} &= -\frac{m}{\sigma} + \frac{1}{\sigma^3} \sum_{i=1}^m (y_i - \omega^T x_i)^2 = 0 \implies \sigma_{ML}^2 = \frac{1}{m} \sum_{i=1}^m (y_i - \omega^T x_i)^2\end{aligned}$$

- BIG Lesson
  - same as the least squared optimization

$$\begin{aligned}\text{loss function} &= \sum_{i=1}^m (y_i - \omega^T x_i)^2 \\ &= \|Y - X\omega\|_2^2 \\ &= (Y - X\omega)^T (Y - X\omega) \\ &= Y^T Y - \omega^T X^T Y - Y^T X\omega + \omega^T X^T X\omega\end{aligned}$$

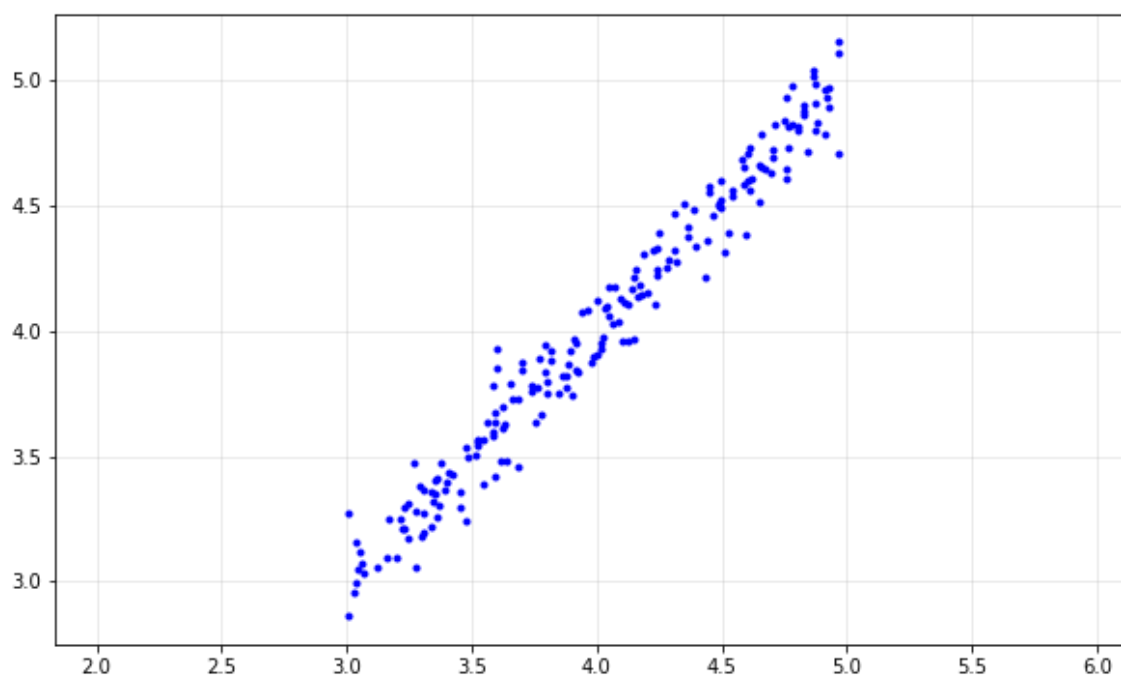
In [7]:

```
m = 200

a = 1
x = 3 + 2*np.random.uniform(0,1,[m,1])
noise = 0.1*np.random.randn(m,1)

y = a*x + noise;
y = np.asmatrix(y)

plt.figure(figsize=(10, 6))
plt.plot(x, y, 'b.')
plt.axis('equal')
plt.grid(alpha=0.3)
plt.show()
```



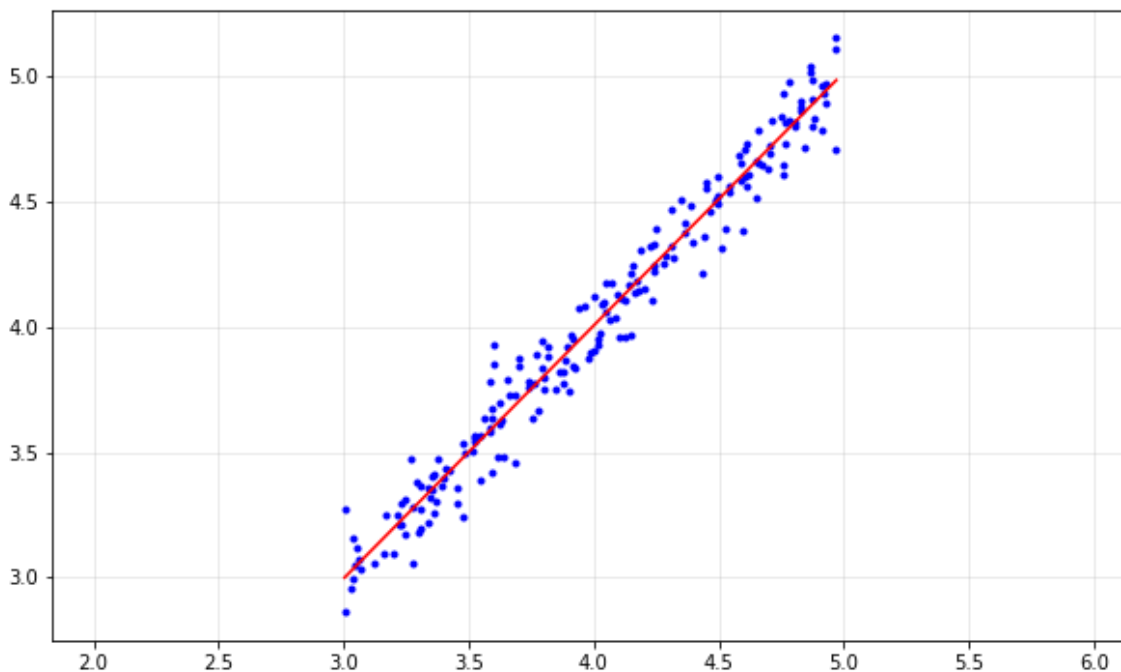
In [8]:

```
# compute theta(1) and theta(2) which are coefficients of  $y = \theta(1)x + \theta(2)$ 
A = np.hstack([np.ones([m, 1]), x])
A = np.asmatrix(A)

theta = (A.T*A).I*A.T*y

# to plot the fitted line
xp = np.linspace(np.min(x), np.max(x))
yp = theta[1,0]*xp + theta[0,0]

plt.figure(figsize=(10, 6))
plt.plot(x, y, 'b.')
plt.plot(xp, yp, 'r')
plt.axis('equal')
plt.grid(alpha=0.3)
plt.show()
```



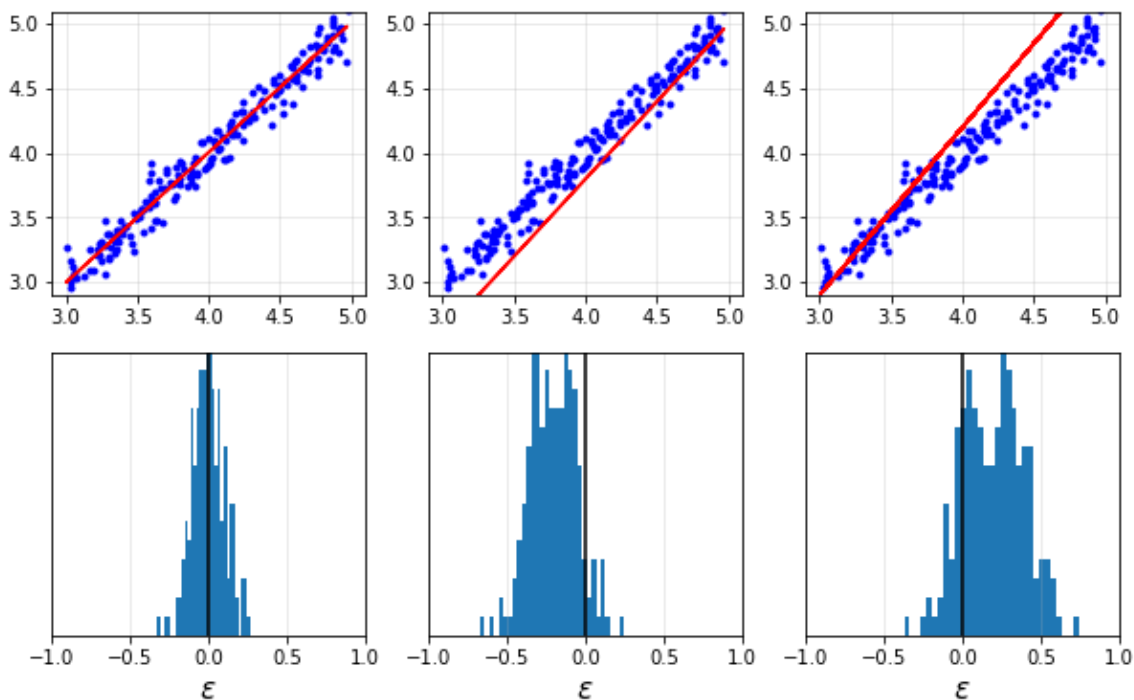
In [9]:

```
yhat0 = theta[1,0]*x + theta[0,0]
err0 = yhat0 - y

yhat1 = 1.2*x - 1
err1 = yhat1 - y

yhat2 = 1.3*x - 1
err2 = yhat2 - y

plt.figure(figsize=(10, 6))
plt.subplot(2,3,1), plt.plot(x,y,'b.',x,yhat0,'r'), plt.axis([2.9, 5.1, 2.9, 5.1]),
plt.grid(alpha=0.3)
plt.subplot(2,3,2), plt.plot(x,y,'b.',x,yhat1,'r'), plt.axis([2.9, 5.1, 2.9, 5.1]),
plt.grid(alpha=0.3)
plt.subplot(2,3,3), plt.plot(x,y,'b.',x,yhat2,'r'), plt.axis([2.9, 5.1, 2.9, 5.1]),
plt.grid(alpha=0.3)
plt.subplot(2,3,4), plt.hist(err0,31), plt.axvline(0, color='k'), plt.xlabel(r'\epsilon',
    fontsize=15),
plt.yticks([]), plt.axis([-1, 1, 0, 15]), plt.grid(alpha=0.3)
plt.subplot(2,3,5), plt.hist(err1,31), plt.axvline(0, color='k'), plt.xlabel(r'\epsilon',
    fontsize=15),
plt.yticks([]), plt.axis([-1, 1, 0, 15]), plt.grid(alpha=0.3)
plt.subplot(2,3,6), plt.hist(err2,31), plt.axvline(0, color='k'), plt.xlabel(r'\epsilon',
    fontsize=15),
plt.yticks([]), plt.axis([-1, 1, 0, 15]), plt.grid(alpha=0.3)
plt.show()
```



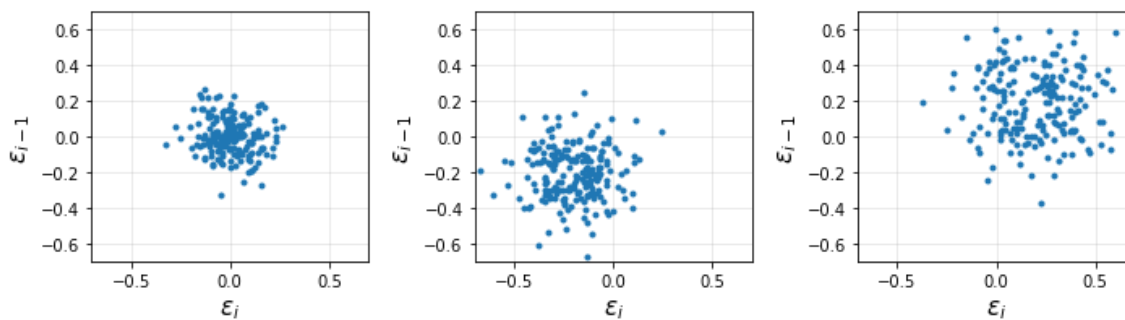
In [10]:

```
a0x = err0[1:]
a0y = err0[0:-1]

a1x = err1[1:]
a1y = err1[0:-1]

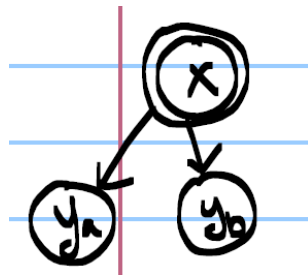
a2x = err2[1:]
a2y = err2[0:-1]

plt.figure(figsize=(10, 3))
plt.subplot(1, 3, 1), plt.plot(a0x, a0y, '.'), plt.axis([-0.7, 0.7, -0.7, 0.7]), plt.grid(alpha=0.3)
plt.xlabel(r'$\epsilon_i$', fontsize=15), plt.ylabel(r'$\epsilon_{i-1}$', fontsize=15)
plt.subplot(1, 3, 2), plt.plot(a1x, a1y, '.'), plt.axis([-0.7, 0.7, -0.7, 0.7]), plt.grid(alpha=0.3)
plt.xlabel(r'$\epsilon_i$', fontsize=15), plt.ylabel(r'$\epsilon_{i-1}$', fontsize=15)
plt.subplot(1, 3, 3), plt.plot(a2x, a2y, '.'), plt.axis([-0.7, 0.7, -0.7, 0.7]), plt.grid(alpha=0.3)
plt.xlabel(r'$\epsilon_i$', fontsize=15), plt.ylabel(r'$\epsilon_{i-1}$', fontsize=15)
plt.tight_layout()
plt.show()
```



## 2.4. Data Fusion with Uncertainties

- Learning Theory (Reza Shadmehr, Johns Hopkins University)  
(<http://www.shadmehrlab.org/Courses/learningtheory.html>)
  - youtube link (<https://www.youtube.com/watch?v=52jIBrAcw9Q>)



$$y_a = x + \varepsilon_a, \quad \varepsilon_a \sim \mathcal{N}(0, \sigma_a^2)$$

$$y_b = x + \varepsilon_b, \quad \varepsilon_b \sim \mathcal{N}(0, \sigma_b^2)$$

- in a matrix form

$$y = \begin{bmatrix} y_a \\ y_b \end{bmatrix} = Cx + \varepsilon = \begin{bmatrix} 1 \\ 1 \end{bmatrix} x + \begin{bmatrix} \varepsilon_a \\ \varepsilon_b \end{bmatrix} \quad \varepsilon \sim \mathcal{N}(0, R), \quad R = \begin{bmatrix} \sigma_a^2 & 0 \\ 0 & \sigma_b^2 \end{bmatrix}$$

$$\begin{aligned} P(y | x) &\sim \mathcal{N}(Cx, R) \\ &= \frac{1}{\sqrt{(2\pi)^2 |R|}} \exp\left(-\frac{1}{2}(y - Cx)^T R^{-1} (y - Cx)\right) \end{aligned}$$

- Find  $\hat{x}_{ML}$

$$\ell = -\log 2\pi - \frac{1}{2} \log |R| - \frac{1}{2} \underbrace{(y - Cx)^T R^{-1} (y - Cx)}$$

$$(y - Cx)^T R^{-1} (y - Cx) = y^T R^{-1} y - y^T R^{-1} Cx - x^T C^T R^{-1} y + x^T C^T R^{-1} Cx$$

$$\Rightarrow \frac{d\ell}{dx} = 0 = -2C^T R^{-1} y + 2C^T R^{-1} Cx$$

$$\therefore x_{ML} = (C^T R^{-1} C)^{-1} C^T R^{-1} y$$

- $(C^T R^{-1} C)^{-1} C^T R^{-1}$

$$\begin{aligned} (C^T R^{-1} C) &= \begin{bmatrix} 1 & 1 \end{bmatrix} \begin{bmatrix} \frac{1}{\sigma_a^2} & 0 \\ 0 & \frac{1}{\sigma_b^2} \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \frac{1}{\sigma_a^2} + \frac{1}{\sigma_b^2} \\ C^T R^{-1} &= \begin{bmatrix} 1 & 1 \end{bmatrix} \begin{bmatrix} \frac{1}{\sigma_a^2} & 0 \\ 0 & \frac{1}{\sigma_b^2} \end{bmatrix} = \begin{bmatrix} \frac{1}{\sigma_a^2} & \frac{1}{\sigma_b^2} \end{bmatrix} \end{aligned}$$

$$\begin{aligned}\hat{x}_{ML} &= (C^T R^{-1} C)^{-1} C^T R^{-1} y = \left( \frac{1}{\sigma_a^2} + \frac{1}{\sigma_b^2} \right)^{-1} \begin{bmatrix} \frac{1}{\sigma_a^2} & \frac{1}{\sigma_b^2} \end{bmatrix} \begin{bmatrix} y_a \\ y_b \end{bmatrix} \\ &= \frac{\frac{1}{\sigma_a^2} y_a + \frac{1}{\sigma_b^2} y_b}{\frac{1}{\sigma_a^2} + \frac{1}{\sigma_b^2}}\end{aligned}$$

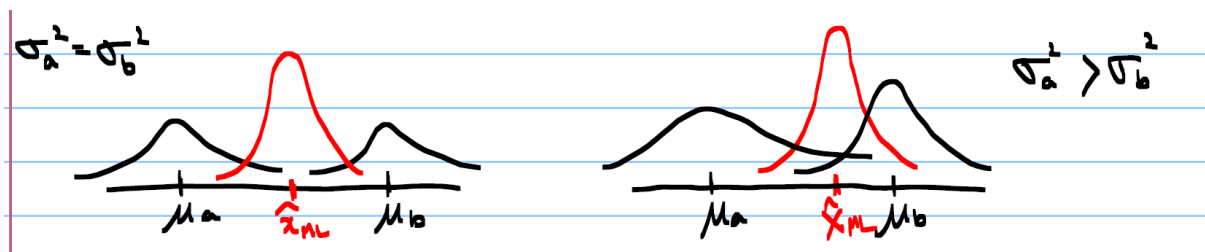
$$\begin{aligned}\text{var}(\hat{x}_{ML}) &= \left( (C^T R^{-1} C)^{-1} C^T R^{-1} \right) \cdot \text{var}(y) \cdot \left( (C^T R^{-1} C)^{-1} C^T R^{-1} \right)^T \\ &= \left( (C^T R^{-1} C)^{-1} C^T R^{-1} \right) \cdot R \cdot \left( (C^T R^{-1} C)^{-1} C^T R^{-1} \right)^T \\ &= (C^T R^{-1} C)^{-1} C^T \cdot (R^{-1})^T C \left( (C^T R^{-1} C)^{-1} \right)^T \\ &= \underbrace{(C^T R^{-1} C)^{-1}} \underbrace{C^T R^{-1} C}_{=I} \left( (C^T R^{-1} C)^{-1} \right)^T = (C^T R^{-1} C)^{-1} \\ &= \frac{1}{\frac{1}{\sigma_a^2} + \frac{1}{\sigma_b^2}} \leq \sigma_a^2, \sigma_b^2\end{aligned}$$

- summary

$$\begin{aligned}\hat{x}_{ML} &= \frac{\frac{1}{\sigma_a^2} y_a + \frac{1}{\sigma_b^2} y_b}{\frac{1}{\sigma_a^2} + \frac{1}{\sigma_b^2}} \\ \text{var}(\hat{x}_{ML}) &= \frac{1}{\frac{1}{\sigma_a^2} + \frac{1}{\sigma_b^2}} \leq \sigma_a^2, \sigma_b^2\end{aligned}$$

- BIG Lesson:

- two sensors are better than one sensor  $\implies$  less uncertainties
- accuracy or uncertainty information is also important in sensors





## Example of two rulers

- 1D example
- how brain works on human measurements from both *haptic* and *visual* channels



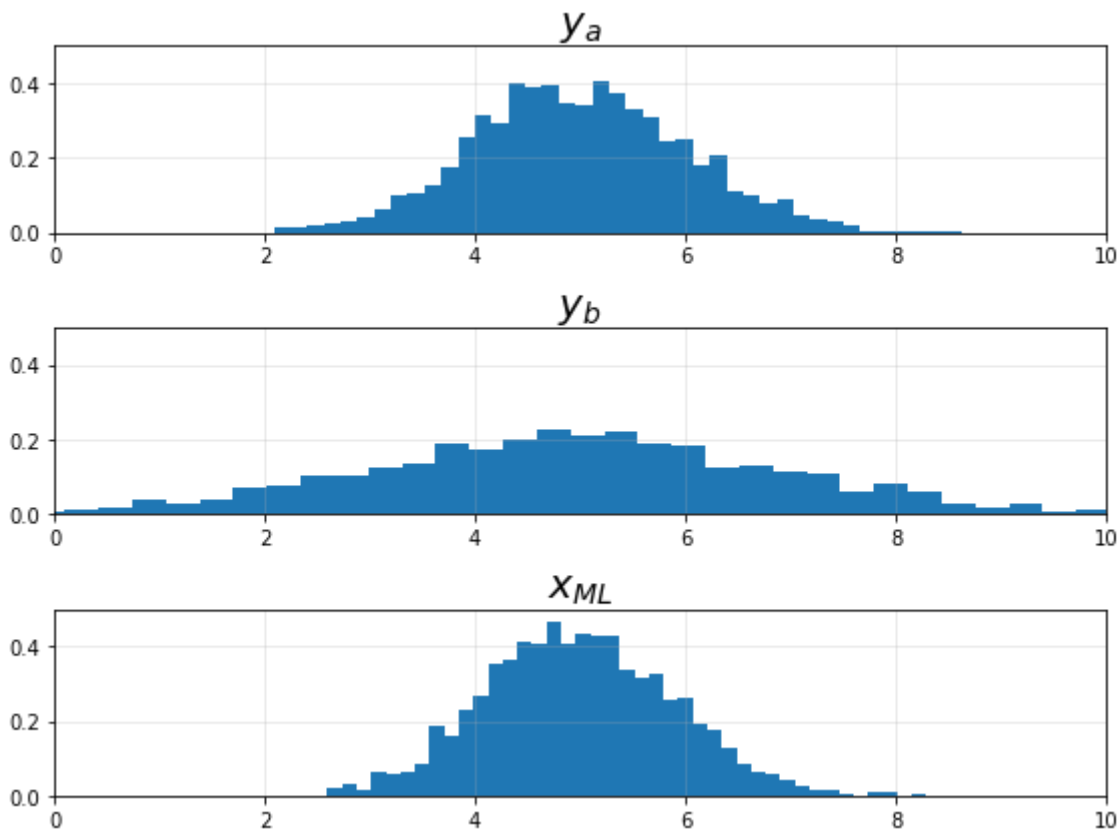
In [11]:

```
x = 5      # true state (length in this example)
a = 1      # sigma of a
b = 2      # sigma of b

YA = []
YB = []
XML = []

for i in range(2000):
    ya = x + np.random.normal(0,a)
    yb = x + np.random.normal(0,b)
    xml = (1/a**2*ya + 1/b**2*yb)/(1/a**2+1/b**2)
    YA.append(ya)
    YB.append(yb)
    XML.append(xml)

plt.figure(figsize=(8, 6))
plt.subplot(3, 1, 1), plt.hist(YA, 41, normed=True), plt.axis([0, 10, 0, 0.5]),
plt.title(r'$y_a$', fontsize=20), plt.grid(alpha=0.3)
plt.subplot(3, 1, 2), plt.hist(YB, 41, normed=True), plt.axis([0, 10, 0, 0.5]),
plt.title(r'$y_b$', fontsize=20), plt.grid(alpha=0.3)
plt.subplot(3, 1, 3), plt.hist(XML, 41, normed=True), plt.axis([0, 10, 0, 0.5]),
plt.title(r'$x_{ML}$', fontsize=20), plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```



## Example of two GPSs

- 2D example

In [12]:

```
x = np.array([5, 10]).reshape(-1, 1) # true position

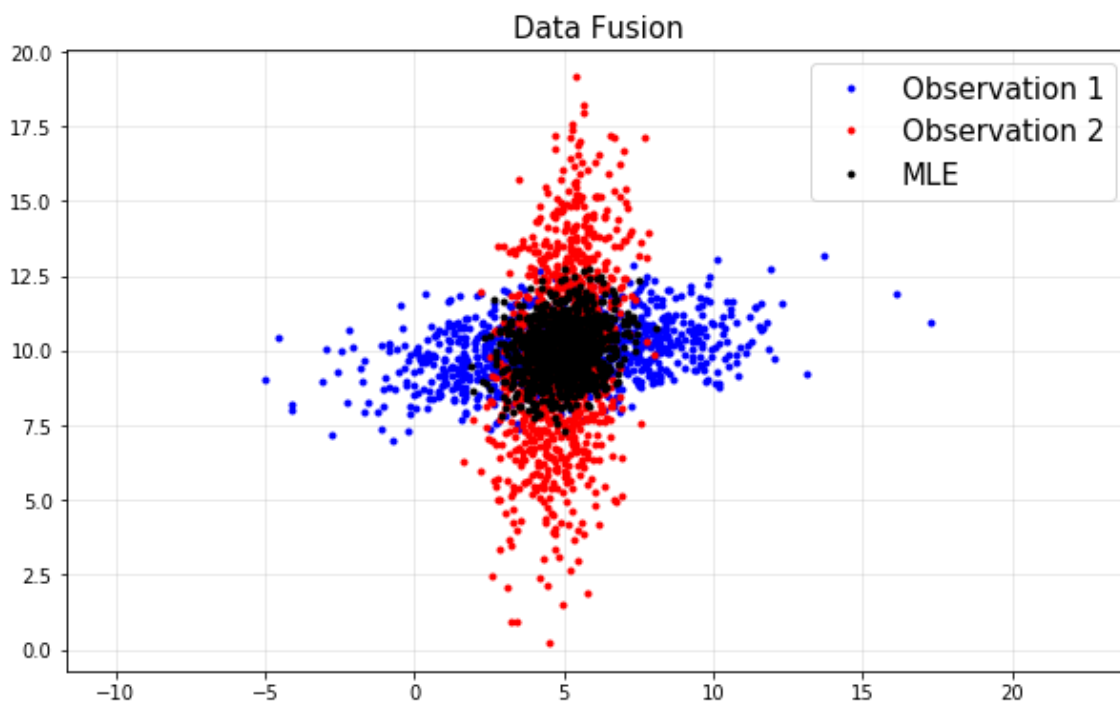
mu = np.array([0, 0])
Ra = np.matrix([[9, 1],
                [1, 1]])
Rb = np.matrix([[1, 1],
                [1, 9]])

YA = []
YB = []
XML = []

for i in range(1000):
    ya = x + np.random.multivariate_normal(mu, Ra).reshape(-1, 1)
    yb = x + np.random.multivariate_normal(mu, Rb).reshape(-1, 1)
    xml = (Ra.I+Rb.I).I*(Ra.I*ya+Rb.I*yb)
    YA.append(ya.T)
    YB.append(yb.T)
    XML.append(xml.T)

YA = np.vstack(YA)
YB = np.vstack(YB)
XML = np.vstack(XML)

plt.figure(figsize=(10, 6))
plt.title('Data Fusion', fontsize=15)
plt.plot(YA[:,0], YA[:,1], 'b.', label='Observation 1')
plt.plot(YB[:,0], YB[:,1], 'r.', label='Observation 2')
plt.plot(XML[:,0], XML[:,1], 'k.', label='MLE')
plt.axis('equal')
plt.grid(alpha=0.3)
plt.legend(fontsize=15)
plt.show()
```



### 3. Maximum-a-Posterior Estimation (MAP)

- Choose  $\theta$  that maximizes the posterior probability of  $\theta$  (i.e. probability in the light of the observed data)
- Posterior probability of  $\theta$  is given by the Bayes Rule

$$P(\theta | D) = \frac{P(D | \theta)P(\theta)}{P(D)}$$

- $P(\theta)$ : Prior probability of  $\theta$  (without having seen any data)
- $P(D | \theta)$ : Likelihood
- $P(D)$ : Probability of the data (independent of  $\theta$ )

$$P(D) = \int P(\theta)P(D | \theta)d\theta$$

- The Bayes rule lets us update our belief about  $\theta$  in the light of observed data
- While doing MAP, we usually maximize the log of the posterior probability

$$\begin{aligned}\theta_{MAP} &= \operatorname{argmax}_{\theta} P(\theta | D) = \operatorname{argmax}_{\theta} \frac{P(D | \theta)P(\theta)}{P(D)} \\ &= \operatorname{argmax}_{\theta} P(D | \theta)P(\theta) \\ &= \operatorname{argmax}_{\theta} \log P(D | \theta)P(\theta) \\ &= \operatorname{argmax}_{\theta} \{\log P(D | \theta) + \log P(\theta)\}\end{aligned}$$

- for multiple observations  $D = \{d_1, d_2, \dots, d_m\}$

$$\theta_{MAP} = \operatorname{argmax}_{\theta} \left\{ \sum_{i=1}^m \log P(d_i | \theta) + \log P(\theta) \right\}$$

- same as MLE except the extra log-prior-distribution term
- MAP allows incorporating our prior knowledge about  $\theta$  in its estimation

$\theta_{MAP} = \operatorname{argmax}_{\theta} P(\theta   D)$	$\theta_{MLE} = \operatorname{argmax}_{\theta} P(D   \theta)$
---	---

### 3.1. MAP for mean of a univariate Gaussian, $\mathcal{N}(\theta, \sigma^2)$

Suppose that  $\theta$  is a random variable with  $\theta \sim \mathcal{N}(\mu, 1^2)$ , but a prior knowledge (unknown  $\theta$  and known  $\mu, \sigma^2$ )

- Observations  $D = \{x_1, x_2, \dots, x_m\}$  : conditionally independent given  $\theta$

$$x_i \sim \mathcal{N}(\theta, \sigma^2)$$

- Joint Probability

$$P(x_1, x_2, \dots, x_m \mid \theta) = \prod_{i=1}^m P(x_i \mid \theta)$$

- MAP: choose  $\theta_{MAP}$

$$\begin{aligned}\theta_{MAP} &= \underset{\theta}{\operatorname{argmax}} P(\theta \mid D) = \frac{P(D \mid \theta)P(\theta)}{P(D)} \\ &= \underset{\theta}{\operatorname{argmax}} P(D \mid \theta)P(\theta) \\ &= \underset{\theta}{\operatorname{argmax}} \{\log P(D \mid \theta) + \log P(\theta)\}\end{aligned}$$

$$\frac{\partial}{\partial \theta}(\log P(D \mid \theta)) = \dots = \frac{1}{\sigma^2} \left( \sum_{i=1}^m x_i - m\theta \right) \quad (\text{we did in MLE})$$

$$\begin{aligned}\frac{\partial}{\partial \theta}(\log P(\theta)) &= \frac{\partial}{\partial \theta} \left( \log \left( \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(\theta-\mu)^2} \right) \right) \\ &\vdots \\ &= \frac{\partial}{\partial \theta} \left( -\frac{1}{2} \log 2\pi - \frac{1}{2}(\theta - \mu)^2 \right) \\ &= \mu - \theta\end{aligned}$$

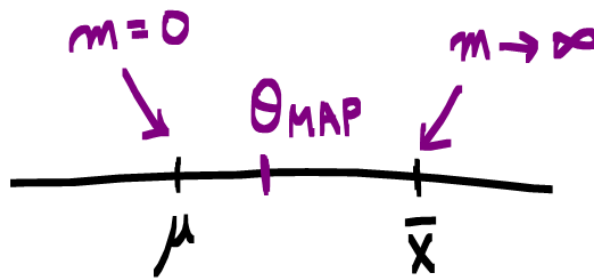
$$\begin{aligned}
&\Rightarrow \frac{\partial}{\partial \theta} (\log P(D | \theta)) + \frac{\partial}{\partial \theta} (\log P(\theta)) \\
&= \frac{1}{\sigma^2} \left( \sum_{i=1}^m x_i - m\theta^* \right) + \mu - \theta^* = 0 \\
&= \frac{1}{\sigma^2} \sum_{i=1}^m x_i + \mu - \left( \frac{m}{\sigma^2} + 1 \right) \theta^* = 0 \\
\theta^* &= \frac{\frac{1}{\sigma^2} \sum_{i=1}^m x_i + \mu}{\frac{m}{\sigma^2} + 1} = \frac{\frac{m}{\sigma^2} \cdot \frac{1}{m} \sum_{i=1}^m x_i + 1 \cdot \mu}{\frac{m}{\sigma^2} + 1}
\end{aligned}$$

$$\therefore \theta_{MAP} = \frac{\frac{m}{\sigma^2}}{\frac{m}{\sigma^2} + 1} \bar{x} + \frac{1}{\frac{m}{\sigma^2} + 1} \mu \quad : \text{look familiar ?}$$

- ML interpretation:

$$\begin{cases} \mu = \text{prior mean} \\ \bar{x} = \text{sample mean} \\ \mu = \text{1st observation} \sim \mathcal{N}(0, 1^2) \\ \bar{x} = \text{2nd observation} \sim \mathcal{N}\left(0, \left(\frac{\sigma}{\sqrt{m}}\right)^2\right) \end{cases}$$

- BIG Lesson: a prior acts as a data



**Note:** prior knowledge

- Education
- Get older
- School ranking

Example) Experiment in class

- Which one do you think is heavier?
  - with eyes closed
  - with visual inspection
  - with haptic (touch) inspection



## 3.2. MAP Python code

Suppose that  $\theta$  is a random variable with  $\theta \sim \mathcal{N}(\mu, 1^2)$ , but a prior knowledge (unknown  $\theta$  and known  $\mu, \sigma^2$ )

$$x_i \sim \mathcal{N}(\theta, \sigma^2)$$

- for mean of a univariate Gaussian

In [13]:

```
# known
mu = 5
sigma = 2

# unknown theta
theta = np.random.normal(mu,1)
x = np.random.normal(theta, sigma)

print('theta = {:.4f}'.format(theta))
print('x = {:.4f}'.format(x))
```

```
theta = 3.8211
x = 5.7443
```

$$\theta_{MAP} = \frac{\frac{m}{\sigma^2}}{\frac{m}{\sigma^2} + 1} \bar{x} + \frac{1}{\frac{m}{\sigma^2} + 1} \mu$$

In [14]:

```
# MAP

m = 4
X = np.random.normal(theta,sigma,[m,1])

xbar = np.mean(X)
theta_MAP = m/(m+sigma**2)*xbar + sigma**2/(m+sigma**2)*mu

print('mu = 5')
print('xbar = {:.4f}'.format(xbar))
print('theta_MAP = {:.4f}'.format(theta_MAP))

mu = 5
xbar = 2.2625
theta_MAP = 3.6313
```

In [15]:

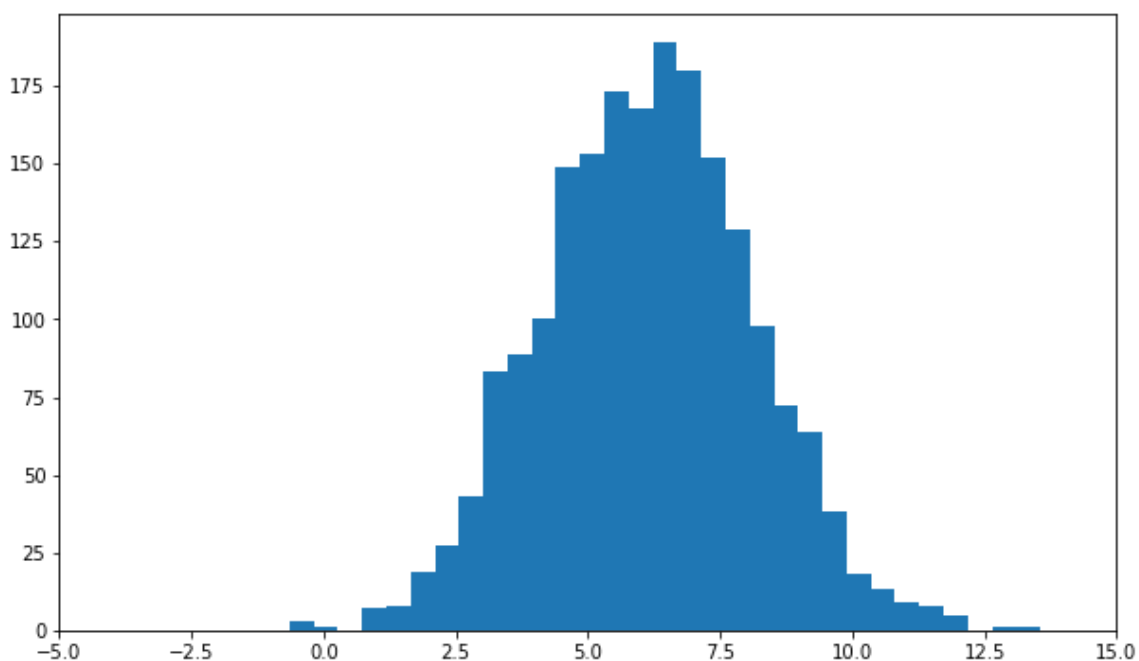
```
# theta
mu = 5
theta = np.random.normal(mu,1)

sigma = 2
m = 2000

X = np.random.normal(theta,sigma,[m,1])
X = np.asmatrix(X)

print('theta = {:.4f}'.format(theta))
plt.figure(figsize=(10,6))
plt.hist(X,31)
plt.xlim([-5,15])
plt.show()
```

theta = 6.1839





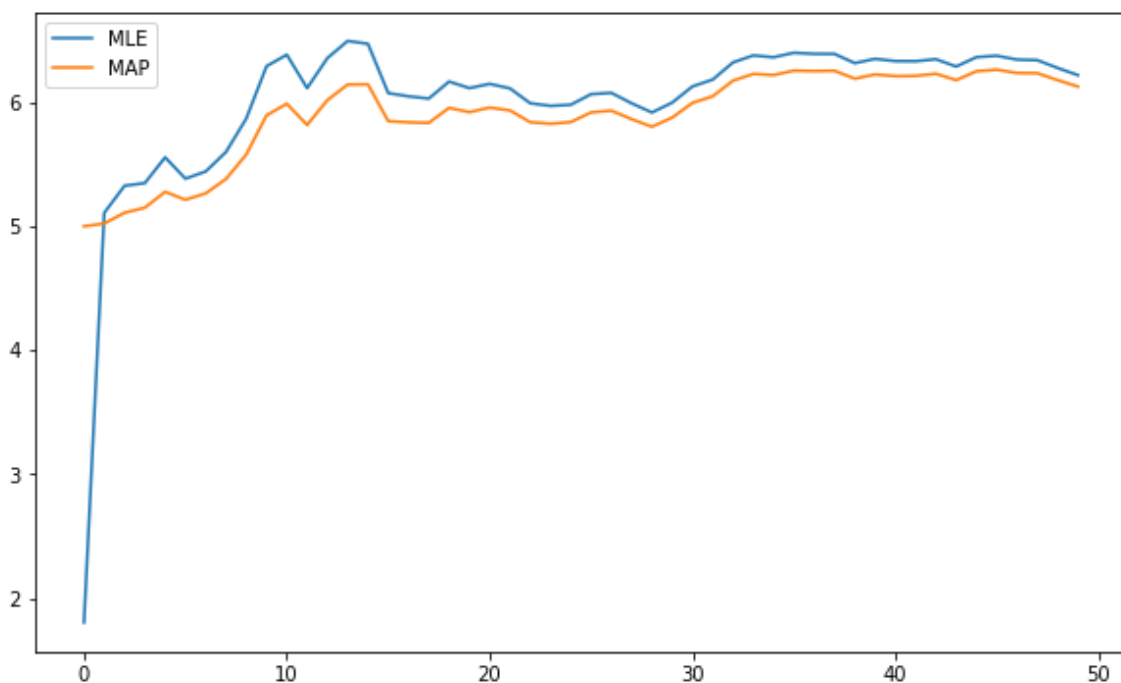
In [16]:

```
n = 50
XMLE = []
XMAP = []

for k in range(n):
    xmle = np.mean(X[0:k+1,0])
    xmap = k/(k+sigma**2)*xmle + sigma**2/(k + sigma**2)*mu
    XMLE.append(xmle)
    XMAP.append(xmap)

print('theta = {:.4f}'.format(theta))
plt.figure(figsize=(10,6))
plt.plot(XMLE)
plt.plot(XMAP)
plt.legend(['MLE', 'MAP'])
plt.show()
```

theta = 6.1839



### 3.3. Object Tracking in Computer Vision

- Optional
- Lecture: Introduction to Computer Vision by Prof. Aaron Bobick at Georgia Tech

In [17]:

```
%%html
<center><iframe src="https://www.youtube.com/embed/rf3DKqWajwY?rel=0"
width="560" height="315" frameborder="0" allowfullscreen></iframe></center>
```

## Tracking as Inference



In [18]:

```
%%html
<center><iframe src="https://www.youtube.com/embed/5yUjYCkm2jI?rel=0"
width="560" height="315" frameborder="0" allowfullscreen></iframe></center>
```

## Tracking with Kalman Filters



# 4. Kernel Density Estimation

- *non-parametric* estimate of density
- Lecture: Learning Theory (Reza Shadmehr, Johns Hopkins University)

In [19]:

```
%%html
```

```
<center><iframe src="https://www.youtube.com/embed/a357NoXy4Nk?rel=0"
width="560" height="315" frameborder="0" allowfullscreen></iframe></center>
```

Kernel density estimation



In [20]:

```
m = 10
mu = 0
sigma = 5

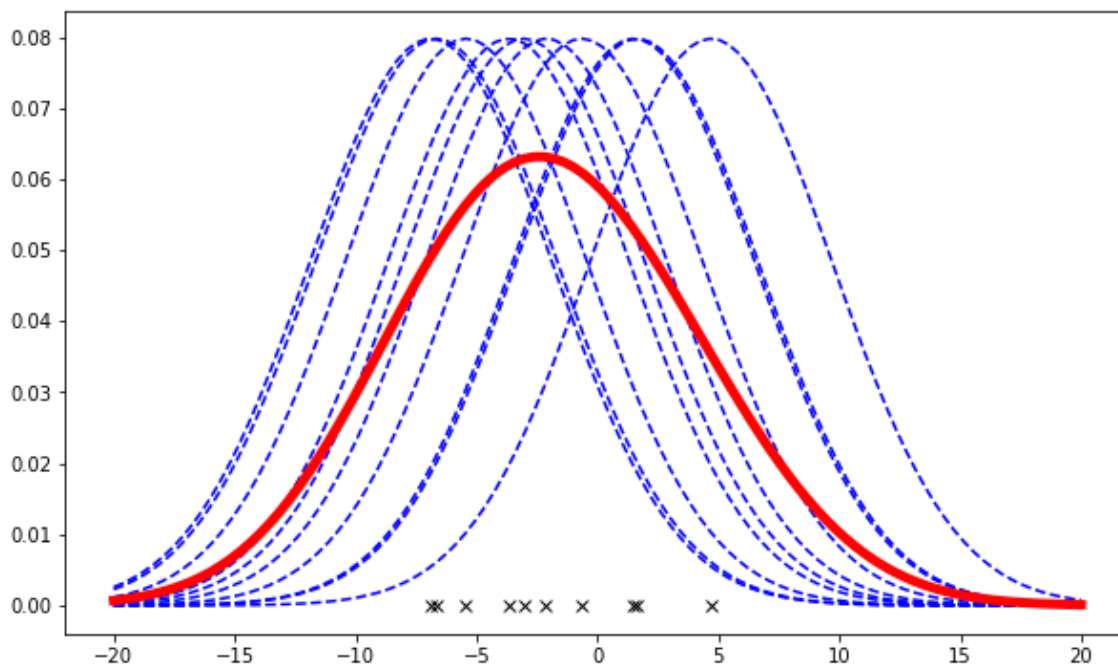
x = np.random.normal(mu,sigma,[m,1])
xp = np.linspace(-20,20,100)
y0 = np.zeros([m,1])

X = []

for i in range(m):
    X.append(norm.pdf(xp,x[i,0],sigma))

X = np.array(X).T
Xnorm = np.sum(X,1)/m

plt.figure(figsize=(10,6))
plt.plot(x,y0,'kx')
plt.plot(xp,X,'b--')
plt.plot(xp,Xnorm,'r',linewidth=5)
plt.show()
```



In [21]:

```
%%javascript
$.getScript('https://kmahelona.github.io/ipython_notebook_goodies/ipython_notebook_toc.js')
```