

# **Autoencoder**


**iSystems Design Lab.**

# Unsupervised Learning


- Definition
  - Unsupervised learning refers to most attempts to extract information from a distribution that do not require human labor to annotate example
  - Main task is to find the 'best' representation of the data
- Dimension Reduction
  - Attempt to compress as much information as possible in a smaller representation
  - Preserve as much information as possible while obeying some constraint aimed at keeping the representation simpler

# Dimension Reduction


- Principal Component Analysis (PCA)
  - Dim reduction without losing too much information




# Dimension Reduction



# Dimension Reduction



# Dimension Reduction




# Autoencoders

- It is like 'deep learning version' of unsupervised learning (dim reduction)
- Definition
  - An autoencoder is a neural network that is trained to attempt to copy its input to its output
  - The network consists of two parts: an encoder and a decoder that produce a reconstruction
- Encoder and Decoder
  - Encoder function :  $h = f(x)$
  - Decoder function :  $x = g(h)$
  - We learn to set  $g(f(x)) = x$

# Autoencoder


- Dimension reduction
- Recover the input data






# Autoencoder

- Dimension reduction
- Recover the input data

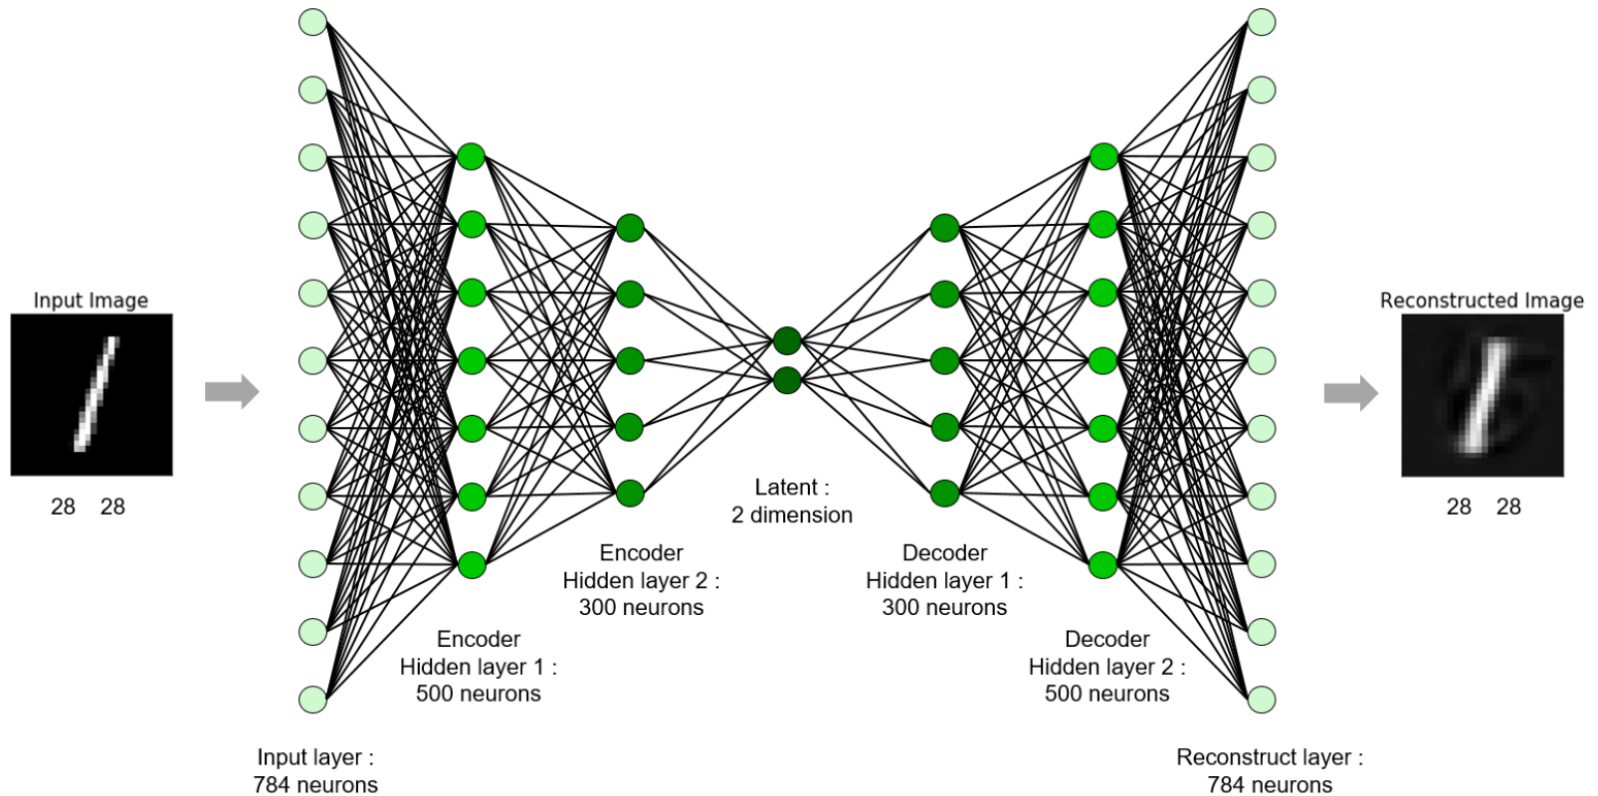


# Autoencoder with TensorFlow

- MNIST example
- Use only (1, 5, 6) digits to visualize in 2-D



# Autoencoder with TensorFlow



# Import Library

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
```

# Load MNIST Data

```
def batch_maker(batch_size, img, label):  
    img_len = len(img)  
    random_idx = np.random.randint(img_len, size = batch_size)  
    return img[random_idx], label[random_idx]
```

```
from tensorflow.examples.tutorials.mnist import input_data  
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

```
Extracting MNIST_data/train-images-idx3-ubyte.gz  
Extracting MNIST_data/train-labels-idx1-ubyte.gz  
Extracting MNIST_data/t10k-images-idx3-ubyte.gz  
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz
```

# Load MNIST Data


```
train_idx = ((np.argmax(mnist.train.labels, 1) == 1) | \
             (np.argmax(mnist.train.labels, 1) == 5) | \
             (np.argmax(mnist.train.labels, 1) == 6))
test_idx = ((np.argmax(mnist.test.labels, 1) == 1) | \
            (np.argmax(mnist.test.labels, 1) == 5) | \
            (np.argmax(mnist.test.labels, 1) == 6))

train_imgs = mnist.train.images[train_idx]
train_labels = mnist.train.labels[train_idx]
test_imgs = mnist.test.images[test_idx]
test_labels = mnist.test.labels[test_idx]
n_train = train_imgs.shape[0]
n_test = test_imgs.shape[0]

print ("Packages loaded")
print ("The number of trainimgs : {}, shape : {}".format(n_train, train_imgs.shape))
print ("The number of testimgs : {}, shape : {}".format(n_test, test_imgs.shape))
```

```
Packages loaded
The number of trainimgs : 16583, shape : (16583, 784)
The number of testimgs : 2985, shape : (2985, 784)
```

# Autoencoder Structure



```
# Shape of input and latent variable
```

```
n_input = 28*28
```

```
# Encoder shape
```

```
n_encoder1 = 500
```

```
n_encoder2 = 300
```

```
n_latent = 2
```

```
# Decoder shape
```

```
n_decoder1 = 300
```

```
n_decoder2 = 500
```

# Weights and Biases

- Define weights and biases for encoder and decoder, separately
- Based on the pre-defined layer size
- Initialize with normal distribution with  $\mu = 0$  and  $\sigma = 0.01$

```
weights = {
    'encoder1' : tf.Variable(tf.random_normal([n_input, n_encoder1], stddev=0.1)),
    'encoder2' : tf.Variable(tf.random_normal([n_encoder1, n_encoder2], stddev=0.1)),
    'latent' : tf.Variable(tf.random_normal([n_encoder2, n_latent], stddev=0.1)),
    'decoder1' : tf.Variable(tf.random_normal([n_latent, n_decoder1], stddev=0.1)),
    'decoder2' : tf.Variable(tf.random_normal([n_decoder1, n_decoder2], stddev=0.1)),
    'reconst' : tf.Variable(tf.random_normal([n_decoder2, n_input], stddev=0.1))
}

biases = {
    'encoder1' : tf.Variable(tf.random_normal([n_encoder1], stddev=0.1)),
    'encoder2' : tf.Variable(tf.random_normal([n_encoder2], stddev=0.1)),
    'latent' : tf.Variable(tf.random_normal([n_latent], stddev=0.1)),
    'decoder1' : tf.Variable(tf.random_normal([n_decoder1], stddev=0.1)),
    'decoder2' : tf.Variable(tf.random_normal([n_decoder2], stddev=0.1)),
    'reconst' : tf.Variable(tf.random_normal([n_input], stddev=0.1))
}

x = tf.placeholder(tf.float32, [None, n_input])
```



# Build a Model

```
def encoder(x, weights, biases):  
    encoder1 = tf.add(tf.matmul(x, weights['encoder1']), biases['encoder1'])  
    encoder1 = tf.nn.tanh(encoder1)  
  
    encoder2 = tf.add(tf.matmul(encoder1, weights['encoder2']), biases['encoder2'])  
    encoder2 = tf.nn.tanh(encoder2)  
  
    latent = tf.add(tf.matmul(encoder2, weights['latent']), biases['latent'])  
  
    return latent
```

```
def decoder(latent, weights, biases):  
    decoder1 = tf.add(tf.matmul(latent, weights['decoder1']), biases['decoder1'])  
    decoder1 = tf.nn.tanh(decoder1)  
  
    decoder2 = tf.add(tf.matmul(decoder1, weights['decoder2']), biases['decoder2'])  
    decoder2 = tf.nn.tanh(decoder2)  
  
    reconst = tf.add(tf.matmul(decoder2, weights['reconst']), biases['reconst'])  
  
    return reconst
```

# Loss, Initializer and Optimizer

- Loss

- Squared loss

$$\frac{1}{N} \sum_{i=1}^N (t_i - y_i)^2$$

- Optimizer

- AdamOptimizer: the most popular optimizer

- Initializer

- Initialize all the empty variables

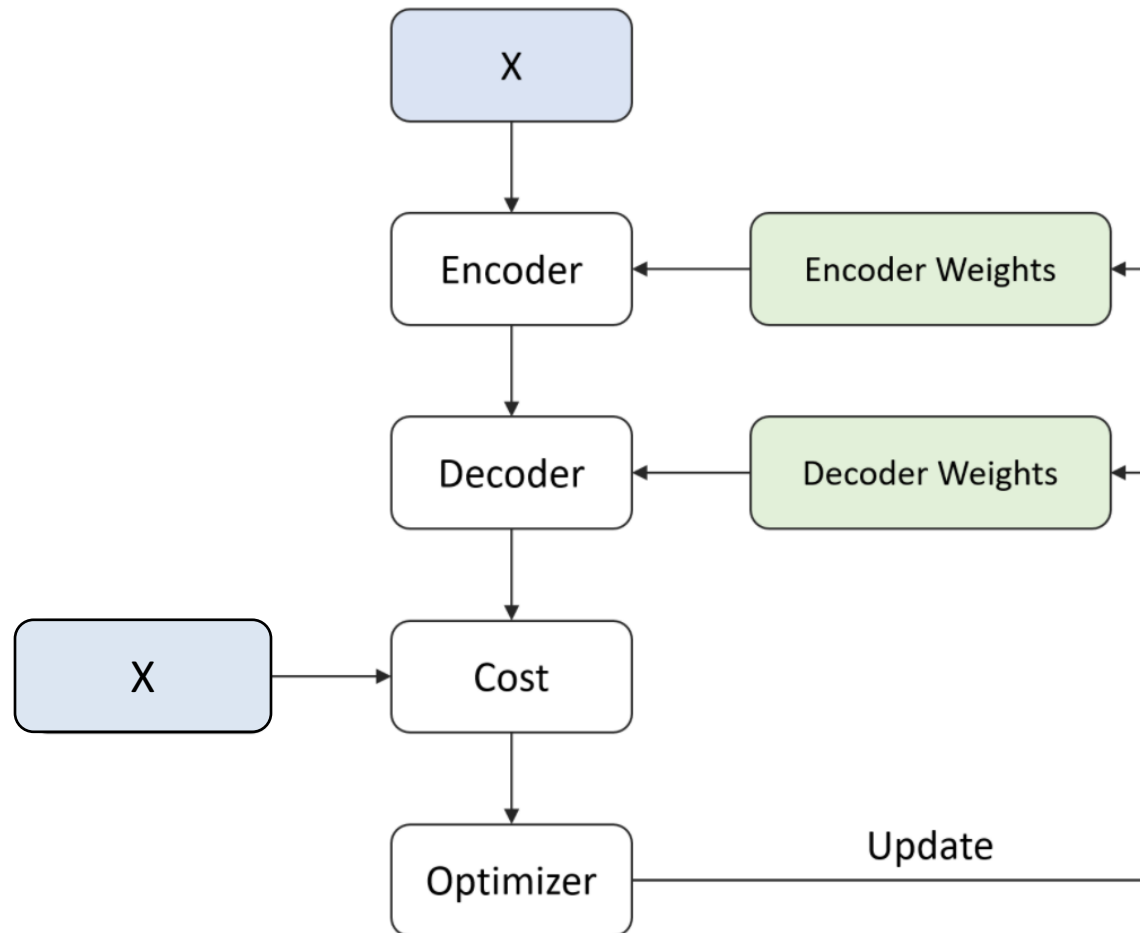
```
LR = 0.0001

latent = encoder(x, weights, biases)
reconst = decoder(latent, weights, biases)
loss = tf.square(tf.subtract(x, reconst))
loss = tf.reduce_mean(loss)

optm = tf.train.AdamOptimizer(LR).minimize(loss)

init = tf.global_variables_initializer()
```

# Summary of Optimization Process



# Iteration Configuration

- Define parameters for training autoencoder
  - n\_batch: batch size for stochastic gradient descent
  - n\_iter: the number of training steps
  - n\_prt: check loss for every n\_prt iteration

```
n_batch = 50  
n_iter = 2500  
n_prt = 250
```

# Optimization

```
# Run initialize
# config = tf.ConfigProto(allow_soft_placement=True) # GPU Allocating policy
# sess = tf.Session(config=config)
sess = tf.Session()
sess.run(init)

# Training cycle
for epoch in range(n_iter):
    train_x, train_y = batch_maker(n_batch, train_imgs, train_labels)
    sess.run(optm, feed_dict={x : train_x})

    if epoch % n_prt == 0:
        c = sess.run(loss, feed_dict={x: train_x})
        print ("Iter : {}".format(epoch))
        print ("Cost : {}".format(c))
```

```
Iter : 0
Cost : 0.4623287618160248
Iter : 250
Cost : 0.04922264814376831
Iter : 500
Cost : 0.040994707494974136
Iter : 750
Cost : 0.04485991969704628
Iter : 1000
Cost : 0.04198655113577843
Iter : 1250
Cost : 0.041495129466056824
Iter : 1500
Cost : 0.04169792681932449
```


# Test or Evaluation

```
test_x, test_y = batch_maker(1, test_imgs, test_labels)
x_reconst = sess.run(reconst, feed_dict={x : test_x})


fig = plt.figure(figsize=(5, 3))
ax1 = fig.add_subplot(1, 2, 1)
ax1.imshow(test_x.reshape(28, 28), 'gray')
ax1.set_title('Input Image', fontsize=15)
ax1.set_xticks([])
ax1.set_yticks([])

ax2 = fig.add_subplot(1, 2, 2)
ax2.imshow(x_reconst.reshape(28, 28), 'gray')
ax2.set_title('Reconstructed Image', fontsize=15)
ax2.set_xticks([])
ax2.set_yticks([])
plt.show()
```

Input Image



Reconstructed Image




# Distribution in Latent Space

- Make a projection of 784-dim image onto 2-dim latent space

```
test_x, test_y = batch_maker(500, test_imgs, test_labels)
test_y = np.argmax(test_y, axis=1)
test_latent = sess.run(latent, feed_dict={x : test_x})

plt.figure(figsize=(10,6))
plt.scatter(test_latent[test_y == 1,0], test_latent[test_y == 1,1], label = 'label = 1')
plt.scatter(test_latent[test_y == 5,0], test_latent[test_y == 5,1], label = 'label = 5')
plt.scatter(test_latent[test_y == 6,0], test_latent[test_y == 6,1], label = 'label = 6')
plt.title('Latent space', fontsize=15)
plt.xlabel('Z1', fontsize=15)
plt.ylabel('Z2', fontsize=15)
plt.legend(fontsize = 15)
plt.show()
```



# Data Generation

```
generate_data = np.array([[-2, 6]])



fig = plt.figure(figsize=(15, 6))
ax = plt.subplot2grid((1, 3), (0, 0), colspan=2)
ax.scatter(test_latent[test_y == 1, 0], test_latent[test_y == 1, 1], label = 'label = 1')
ax.scatter(test_latent[test_y == 5, 0], test_latent[test_y == 5, 1], label = 'label = 5')
ax.scatter(test_latent[test_y == 6, 0], test_latent[test_y == 6, 1], label = 'label = 6')
ax.scatter(generate_data[:, 0], generate_data[:, 1], label = 'generate', s = 150, c = 'k', marker
ax.set_title('Latent space', fontsize=15)
ax.set_xlabel('Z1', fontsize=15)
ax.set_ylabel('Z2', fontsize=15)
ax.legend(fontsize = 15)

latent_input = tf.placeholder(tf.float32, [None, n_latent])
reconst = decoder(latent_input, weights, biases)
generate_x = sess.run(reconst, feed_dict={latent_input : generate_data})

ax = plt.subplot2grid((1, 3), (0, 2), colspan=1)
ax.imshow(generate_x.reshape(28, 28), 'gray')
ax.set_title('Generate Image', fontsize=15)
ax.set_xticks([])
ax.set_yticks([])
plt.show()
```



# Data Generation: Generative Model



# Visualization

```
# Initialize canvas
nx = ny = 20
x_values = np.linspace(-4, 4, nx)
y_values = np.linspace(-4, 10, ny)
canvas = np.empty((28*ny, 28*nx))

# Define placeholder
latent_input = tf.placeholder(tf.float32, [None, n_latent])
reconst = decoder(latent_input, weights, biases)

for i, yi in enumerate(y_values):
    for j, xi in enumerate(x_values):
        latent_ = np.array([[xi, yi]])
        reconst_ = sess.run(reconst, feed_dict={latent_input : latent_})
        canvas[(nx-i-1)*28:(nx-i)*28, j*28:(j+1)*28] = reconst_.reshape(28, 28)

plt.figure(figsize=(10, 10))
plt.imshow(canvas, clim=(0, 1), cmap=plt.cm.jet)
plt.title('Manifold', fontsize=15)
plt.xticks([])
plt.xlabel('Z1', fontsize=15)
plt.yticks([])
plt.ylabel('Z2', fontsize=15)
plt.show()
```

# Visualization

