

Learning from Imbalanced Data


Piyush Rai

Machine Learning (CS771A)

Nov 5, 2016

Learning from Imbalanced Data


- Consider binary classification
- Often the classes are highly imbalanced



- Should I feel happy if my classifier gets 99.997% **classification accuracy** on test data ?

Learning from Imbalanced Data

- Other problems can also exhibit imbalance (e.g., binary matrix completion)



Binary Matrix Completion
0.001 % 1s in the matrix

- Should I feel happy if my matrix completion model gets 99.999% **matrix completion accuracy** (or MAE close to 0) on the test entries?

True Definition of Imbalance Data?

- Debatable..
- Scenario 1: 100,000 negative and 1000 positive examples
- Scenario 2: 10,000 negative and 10 positive examples
- Scenario 3: 1000 negative and 1 positive example
- Usually, imbalance is characterized by absolute rather than relative rarity
 - Finding needles in a haystack..

Minimizing Loss

- Any model to minimize the loss, e.g.,

$$\text{Classification: } \hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \sum_{n=1}^N \ell(y_n, \mathbf{w}^\top \mathbf{x}_n)$$

$$\text{Matrix Completion: } (\hat{\mathbf{U}}, \hat{\mathbf{V}}) = \arg \min_{\mathbf{U}, \mathbf{V}} \|\mathbf{X} - \mathbf{UV}^\top\|^2$$

.. will usually get a high accuracy

- However, it will be highly biased towards predicting the majority class
 - Thus accuracy alone can't be trusted as the evaluation measure if we care more about predicting minority class (say positive) correctly

Better Evaluation Measures

- Precision: What fraction of positive predictions is truly positive

$$P = \frac{\# \text{ example correctly predicted as positive}}{\# \text{ examples predicted as positive}}$$

- Recall: What fraction of total positives are predicted as positives

$$R = \frac{\# \text{ example correctly predicted as positive}}{\# \text{ total positive examples in the test set}}$$

data	label	predicted
------	-------	-----------



0

0



0

1



1

0



1

1



0

1



1

1



0

0

$$\text{precision} = \frac{2}{4}$$

$$\text{recall} = \frac{2}{3}$$

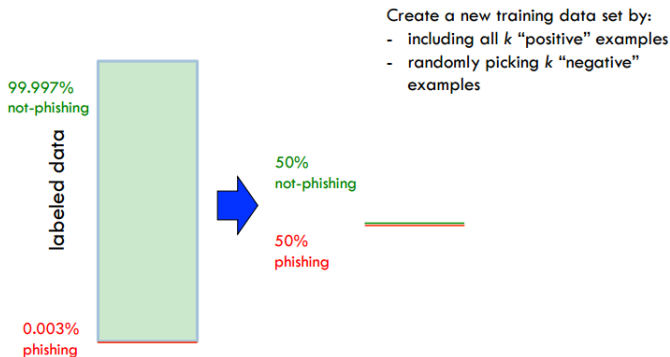
- Often there is a trade-off between precision and recall. Also these can be combined to yield other measures such as F1 score, AUC score, etc.

Dealing with Class Imbalance

- Modifying the training data (the class distribution)
 - Undersampling the majority class
 - Oversampling the minority class
 - Reweighting the examples
- Modifying the learning model
 - Use loss functions customized to handle class imbalance
- Reweighting can be also seen as a way to modify the loss function

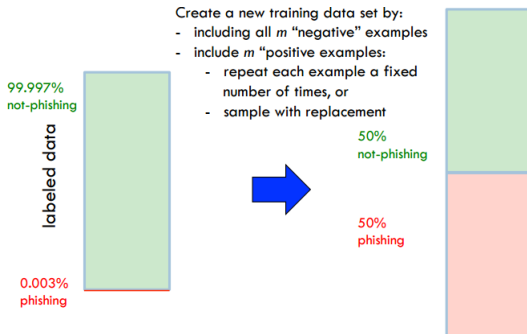
Modifying the Training Data

Undersampling



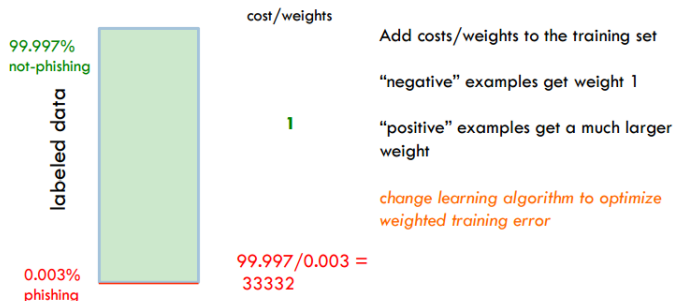
- Throws away a lot of data/information. But efficient to train

Oversampling



- From the loss function's perspective, the repeated examples simply contribute multiple times to the loss function
- Oversampling usually tends to perform undersampling because we are using more data to train the model
- Some oversampling methods (SMOTE) are based on creating synthetic examples from the minority class

Reweighting Examples



- Similar effect as oversampling but is more efficient (because there is no multiplicity of examples)
- Also requires a classifier that can learn with weighted examples

Modifying the Loss Function

Loss Functions Customized for Imbalanced Data

- Traditional loss functions have the form: $\sum_{n=1}^N \ell(y_n, f(\mathbf{x}_n))$
- Such loss functions look at positive and negative examples individually, so the majority class tends to overwhelm the minority class
- Reweighting the loss function differently for different classes can be one way to handle class imbalance, e.g., $\sum_{n=1}^N c_{y_n} \ell(y_n, f(\mathbf{x}_n))$
- Alternatively, we can use loss functions that look at **pairs of examples** (a positive example \mathbf{x}_n^+ and a negative example \mathbf{x}_m^-). For example:

$$\ell(f(\mathbf{x}_n^+), f(\mathbf{x}_m^-)) = \begin{cases} 0, & \text{if } f(\mathbf{x}_n^+) > f(\mathbf{x}_m^-) \\ 1, & \text{otherwise} \end{cases}$$

- These are called **“pairwise” loss functions**
- Why is it a good loss function for imbalanced data?

Pairwise Loss Functions

- Using **pairs** with one +ve and one -ve doesn't let one class overwhelm other

$$\sum_{n=1}^{N_+} \sum_{m=1}^{N_-} \ell(f(\mathbf{x}_n^+), f(\mathbf{x}_m^-)) + \lambda R(f)$$

- The pairwise loss function **only cares about the difference** between scores of a pair of positive and negative examples (which is actually a good thing!)
 - Minimizing the above loss w.r.t. f give us an f that tends to give positive examples a higher score than the negative examples, which is similar in spirit to maximizing the AUC (Area Under the ROC Curve) score
 - AUC (intuitively): The probability that a **randomly chosen** pos. example will have a higher score than a **randomly chosen** neg. example
 - Empirical AUC of f on a training set with N_+ and N_- pos. and neg. ex.

$$\text{AUC}(f) = \frac{1}{N_+ N_-} \sum_{n=1}^{N_+} \sum_{m=1}^{N_-} \mathbb{1}(f(\mathbf{x}_n^+) > f(\mathbf{x}_m^-))$$

- Note: Commonly used pairwise loss functions act as a proxy of the (negative) AUC score (or of closely related measures such as F1 score)

Pairwise Loss Functions

- A proxy based on hinge-loss like pairwise loss function for a linear model

$$\ell(\mathbf{w}, \mathbf{x}_n^+, \mathbf{x}_m^-) = \max\{0, 1 - (\mathbf{w}^\top \mathbf{x}_n^+ - \mathbf{w}^\top \mathbf{x}_m^-)\} = \max\{0, 1 - \mathbf{w}^\top (\mathbf{x}_n^+ - \mathbf{x}_m^-)\}$$

- It basically says that the difference between scores of positive and negative examples should be at least 1 (which is like a “margin”)
- The overall objective will have the form

$$\frac{\|\mathbf{w}\|^2}{2} + \sum_{n=1}^{N_+} \sum_{m=1}^{N_-} \ell(\mathbf{w}, \mathbf{x}_n^+, \mathbf{x}_m^-)$$

- Convex objective (if using the hinge loss). Can be efficiently optimized using stochastic optimization (see “Online AUC Maximization”, Zhao et al, 2011)
- Note: Similar ideas can be used for solving binary matrix factorization and matrix completion problems as well[†]
 - E.g., if matrix entry $X_{nm} = 1$ and $X_{nm'} = -1$ then **loss=0** (or “small”) if $\mathbf{u}_n^\top \mathbf{v}_m > \mathbf{u}_n^\top \mathbf{v}_{m'}$. E.g., $\ell(\mathbf{U}, \mathbf{V}, n, m, m') = -\log \sigma(\mathbf{u}_n^\top \mathbf{v}_m - \mathbf{u}_n^\top \mathbf{v}_{m'})$

[†] If interested, see: “BPR: Bayesian Personalized Ranking from Implicit Feedback” (Rendle et al, 2009)

Summary

- Imbalanced data needs to be handled with care
- Classification accuracies can be very misleading for such data
 - Should look at measures such as precision, recall, or other variants that are robust to class imbalance
- Sampling heuristics work reasonably on many data sets
- More principled approaches are based on **modifying the loss function**
 - Instead of minimizing the classification error, optimize w.r.t. other metrics such as precision, recall, F1 score, AUC, etc.
- Another way to look at this problem could be as an **anomaly detection** problem (minority class is anomaly) or **density estimation** problem