

Supervised Learning

without Scikit Learn


by Prof. Seungchul Lee
Industrial AI Lab
<http://isystems.unist.ac.kr/>
POSTECH


Table of Contents

- I. 1. Optimization
- II. 2. Linear Regression
- III. 3. Classification (Linear)
 - I. 3.1. Distance
 - II. 3.2. Illustrative Example
 - I. 3.2.1. Optimization Formulation 1
 - II. 3.2.2. Outlier
 - III. 3.2.3. Optimization Formulation 2
 - III. 3.3. Maximize Margin (Finally, it is Support Vector Machine)
 - IV. 3.4. Logistic Regression

1. Optimization

- an important tool in 1) engineering problem solving and 2) decision science
- people optimize
- nature optimizes





(source: <http://nautil.us/blog/to-save-drowning-people-ask-yourself-what-would-light-do> (<http://nautil.us/blog/to-save-drowning-people-ask-yourself-what-would-light-do>))

3 key components

1. objective
2. decision variable or unknown
3. constraints

Procedures

1. The process of identifying objective, variables, and constraints for a given problem is known as "modeling"
2. Once the model has been formulated, optimization algorithm can be used to find its solutions.

In mathematical expression

$$\begin{aligned} \min_x \quad & f(x) \\ \text{subject to} \quad & g_i(x) \leq 0, \quad i = 1, \dots, m \end{aligned}$$

Remarks) equivalent

$$\begin{aligned} \min_x f(x) & \Leftrightarrow \max_x -f(x) \\ g_i(x) \leq 0 & \Leftrightarrow -g_i(x) \geq 0 \\ h(x) = 0 & \Leftrightarrow \begin{cases} h(x) \leq 0 \\ h(x) \geq 0 \end{cases} \text{ and} \end{aligned}$$

The good news: for many classes of optimization problems, people have already done all the "hardwork" of developing numerical algorithms

$$\begin{array}{ll}
\max_x & x_1 + x_2 \\
\text{subject to} & 2x_1 + x_2 \leq 29 \\
& x_1 + 2x_2 \leq 25 \\
& x_1 \geq 2 \\
& x_2 \geq 5
\end{array}
\implies
\begin{array}{ll}
\min_x & -[1 \quad 1] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \\
\text{subject to} & \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \leq \begin{bmatrix} 29 \\ 25 \end{bmatrix} \\
& \begin{bmatrix} 2 \\ 5 \end{bmatrix} \leq \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \leq \begin{bmatrix} \\ \end{bmatrix}
\end{array}$$

```

In [1]: import numpy as np
import matplotlib.pyplot as plt
import cvxpy as cvx

f = np.array([[1, 1]])
A = np.array([[2, 1], [1, 2]])
b = np.array([[29], [25]])
lb = np.array([[2], [5]])

x = cvx.Variable(2,1)
obj = cvx.Minimize(-f*x)
const = [A*x <= b, lb <= x]

prob = cvx.Problem(obj, const).solve()

print (x.value)

[[11.]
 [ 7.]]

```

2. Linear Regression

Begin by considering linear regression (easy to extend to more complex predictions later on)

Given $\begin{cases} x_i : \text{inputs} \\ y_i : \text{outputs} \end{cases}$, find θ_1 and θ_2

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix}, \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix} \approx \hat{y}_i = \theta_1 x_i + \theta_2$$


- \hat{y}_i : predicted output

- $\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$: Model parameters

$$\hat{y}_i = f(x_i, \theta) \text{ in general}$$

- In many cases, a linear model to predict y_i can be used


$$\hat{y}_i = \theta_1 x_i + \theta_2 \quad \text{such that} \quad \min_{\theta_1, \theta_2} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$



```
In [2]: import numpy as np
import matplotlib.pyplot as plt

# data points in column vector [input, output]
x = np.array([0.1, 0.4, 0.7, 1.2, 1.3, 1.7, 2.2, 2.8, 3.0, 4.0, 4.3, 4.4, 4.9]).reshape(-1, 1)
y = np.array([0.5, 0.9, 1.1, 1.5, 1.5, 2.0, 2.2, 2.8, 2.7, 3.0, 3.5, 3.7, 3.9]).reshape(-1, 1)

# to plot
plt.figure(figsize=(10, 6))
plt.plot(x, y, 'ko', label="data")
plt.xlabel('X', fontsize=15)
plt.ylabel('Y', fontsize=15)
plt.axis('scaled')
plt.grid(alpha=0.3)
plt.xlim([0, 5])
plt.show()
```



Use CVXPY optimization (least squared)

For convenience, we define a function that maps inputs to feature vectors, ϕ

$$\begin{aligned}\hat{y}_i &= [x_i \quad 1] \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} \\ &= \begin{bmatrix} x_i \\ 1 \end{bmatrix}^T \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}, \quad \text{feature vector } \phi(x_i) = \begin{bmatrix} x_i \\ 1 \end{bmatrix} \\ &= \phi^T(x_i)\theta\end{aligned}$$

$$\Phi = \begin{bmatrix} x_1 & 1 \\ x_2 & 1 \\ \vdots & \\ x_m & 1 \end{bmatrix} = \begin{bmatrix} \phi^T(x_1) \\ \phi^T(x_2) \\ \vdots \\ \phi^T(x_m) \end{bmatrix} \implies \hat{y} = \begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_m \end{bmatrix} = \Phi\theta$$

Model parameter estimation

$$\min_{\theta} \|\hat{y} - y\|_2 = \min_{\theta} \|\Phi\theta - y\|_2$$

```
In [3]: import cvxpy as cvx

m = y.shape[0]
#A = np.hstack([x, np.ones([m, 1])])
A = np.hstack([x, x**0])
A = np.asmatrix(A)

theta2 = cvx.Variable(2, 1)
obj = cvx.Minimize(cvx.norm(A*theta2-y, 2))
cvx.Problem(obj, []).solve()


print('theta:\n', theta2.value)

theta:
[[0.67129519]
 [0.65306531]]
```

```
In [4]: # to plot
plt.figure(figsize=(10, 6))
plt.title('$L_2$ Regression', fontsize=15)
plt.xlabel('X', fontsize=15)
plt.ylabel('Y', fontsize=15)
plt.plot(x, y, 'ko', label="data")

# to plot a straight line (fitted line)
xp = np.arange(0, 5, 0.01).reshape(-1, 1)
theta2 = theta2.value
yp = theta2[0,0]*xp + theta2[1,0]


plt.plot(xp, yp, 'r', linewidth=2, label="$L_2$")
plt.legend(fontsize=15)
plt.axis('scaled')
plt.grid(alpha=0.3)
plt.xlim([0, 5])
plt.show()
```



3. Classification (Linear)

- Figure out, autonomously, which category (or class) an unknown item should be categorized into
- Number of categories / classes
 - Binary: 2 different classes
 - Multiclass: more than 2 classes
- Feature
 - The measurable parts that make up the unknown item (or the information you have available to categorize)

- Perceptron: make use of sign of data
 - Discuss it later
- Logistic regression is a classification algorithm
 - don't be confused
- To find a classification boundary




Sign


- Sign with respect to a line

$$\omega = \begin{bmatrix} \omega_1 \\ \omega_2 \end{bmatrix}, \quad x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \implies g(x) = \omega_1 x_1 + \omega_2 x_2 + \omega_0 = \omega^T x + \omega_0$$


$$\omega = \begin{bmatrix} \omega_0 \\ \omega_1 \\ \omega_2 \end{bmatrix}, \quad x = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} \implies g(x) = \omega_0 + \omega_1 x_1 + \omega_2 x_2 = \omega^T x$$



Perceptron




- Hyperplane
 - Separates a D-dimensional space into two half-spaces
 - Defined by an outward pointing normal vector
 - ω is orthogonal to any vector lying on the hyperplane



How to find ω

- All data in class 1
 - $g(\omega^T x) > 0$
- All data in class 0
 - $g(\omega^T x) < 0$



Perceptron Algorithm

The perceptron implements

$$h(x) = \text{sign}(\omega^T x)$$

Given the training set


$$(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N) \quad \text{where } y_i \in \{-1, 1\}$$

1) pick a misclassified point

$$\text{sign}(\omega^T x_n) \neq y_n$$

2) and update the weight vector

$$\omega \leftarrow \omega + y_n x_n$$



- Why perceptron updates work ?
- Let's look at a misclassified positive example ($y_n = +1$)
 - perceptron (wrongly) thinks $\omega_{old}^T x_n < 0$
- updates would be

$$\omega_{new} = \omega_{old} + y_n x_n = \omega_{old} + x_n$$

$$\omega_{new}^T x_n = (\omega_{old} + x_n)^T x_n = \omega_{old}^T x_n + x_n^T x_n$$

- Thus $\omega_{new}^T x_n$ is less negative than $\omega_{old}^T x_n$

```
In [5]: import numpy as np
import matplotlib.pyplot as plt

% matplotlib inline
```

```
In [6]: #training data gereneration
m = 100
x1 = 8*np.random.rand(m, 1)
x2 = 7*np.random.rand(m, 1) - 4

g0 = 0.8*x1 + x2 - 3
g1 = g0 - 1
g2 = g0 + 1
```


```
In [7]: C1 = np.where(g1 >= 0)
C2 = np.where(g2 < 0)
print(C1)

(array([ 0,  2,  4,  5,  6,  7, 10, 17, 22, 24, 25, 27, 28, 30, 31, 38, 51,
        52, 53, 56, 57, 61, 62, 64, 70, 78, 85, 86, 89, 98], dtype=int64), a
rray([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0], dtype=int64))
```

```
In [8]: C1 = np.where(g1 >= 0)[0]
C2 = np.where(g2 < 0)[0]
print(C1.shape)
print(C2.shape)

(30,)
(41,)
```

```
In [9]: plt.figure(figsize=(10, 6))
plt.plot(x1[C1], x2[C1], 'ro', label='C1')
plt.plot(x1[C2], x2[C2], 'bo', label='C2')
plt.title('Linearly seperable classes', fontsize=15)
plt.legend(loc='upper left', fontsize=15)
plt.xlabel(r'$x_1$', fontsize=20)
plt.ylabel(r'$x_2$', fontsize=20)
plt.show()
```



$$x = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ (x^{(3)})^T \\ \vdots \\ (x^{(m)})^T \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} \\ 1 & x_1^{(3)} & x_2^{(3)} \\ \vdots & \vdots & \vdots \\ 1 & x_1^{(m)} & x_2^{(m)} \end{bmatrix}$$

$$y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ y^{(3)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

```
In [10]: X1 = np.hstack([np.ones([C1.shape[0],1]), x1[C1], x2[C1]])
X2 = np.hstack([np.ones([C2.shape[0],1]), x1[C2], x2[C2]])
X = np.vstack([X1, X2])

y = np.vstack([np.ones([C1.shape[0],1]), -np.ones([C2.shape[0],1])])

X = np.asmatrix(X)
y = np.asmatrix(y)
```

$$\omega = \begin{bmatrix} \omega_1 \\ \omega_2 \\ \omega_3 \end{bmatrix}$$

$$\omega \leftarrow \omega + yx$$

where (x, y) is a misclassified training point

```
In [11]: w = np.ones([3,1])
w = np.asmatrix(w)

n_iter = y.shape[0]
for k in range(n_iter):
    for i in range(n_iter):
        if y[i,0] != np.sign(X[i,:]*w)[0,0]:
            w += y[i,0]*X[i,:].T

print(w)

[[-13.          ]
 [  3.35733863]
 [  8.90909811]]
```


$$g(x) = \omega^T x + \omega_0 = \omega_1 x_1 + \omega_2 x_2 + \omega_0 = 0$$

$$\implies x_2 = -\frac{\omega_1}{\omega_2} x_1 - \frac{\omega_0}{\omega_2}$$


Not a unique solution

```
In [12]: x1p = np.linspace(0,8,100).reshape(-1,1)
x2p = - w[1,0]/w[2,0]*x1p - w[0,0]/w[2,0]

plt.figure(figsize=(10, 6))
plt.scatter(x1[C1], x2[C1], c='r', s=50, label='C1')
plt.scatter(x1[C2], x2[C2], c='b', s=50, label='C2')
plt.plot(x1p, x2p, c='k', label='perceptron')
plt.xlim([0,8])
plt.xlabel('$x_1$', fontsize = 20)
plt.ylabel('$x_2$', fontsize = 20)
plt.legend(loc = 4, fontsize = 15)
plt.show()
```




Perceptron



3.1. Distance

$$\omega = \begin{bmatrix} \omega_1 \\ \omega_2 \end{bmatrix}, x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \implies g(x) = \omega^T x + \omega_0 = \omega_1 x_1 + \omega_2 x_2 + \omega_0$$



- Find a distance between $g(x) = -1$ and $g(x) = 1$

suppose $g(x_1) = -1$, $g(x_2) = 1$

$$\begin{aligned} \omega^T x_1 + \omega_0 &= -1 \\ \omega^T x_2 + \omega_0 &= 1 \end{aligned} \implies \omega^T (x_2 - x_1) = 2$$

$$s = \left\langle \frac{\omega}{\|\omega\|}, x_2 - x_1 \right\rangle = \frac{1}{\|\omega\|} \omega^T (x_2 - x_1) = \frac{2}{\|\omega\|}$$

3.2. Illustrative Example

- Binary classification: C_1 and C_2
- Features: the coordinate of i th data

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

- Is it possible to distinguish between C_1 and C_2 by its coordinates?
- We need to find a separating hyperplane (or a line in 2D)

$$\begin{aligned} \omega_1 x_1 + \omega_2 x_2 + \omega_0 &= 0 \\ \begin{bmatrix} \omega_1 & \omega_2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + \omega_0 &= 0 \\ \omega^T x + \omega_0 &= 0 \end{aligned}$$

```
In [13]: import numpy as np
import matplotlib.pyplot as plt

#training data generation
x1 = 8*np.random.rand(100, 1)
x2 = 7*np.random.rand(100, 1) - 4

g0 = 0.8*x1 + x2 - 3
g1 = g0 - 1
g2 = g0 + 1


C1 = np.where(g1 >= 0)[0]
C2 = np.where(g2 < 0)[0]
```

```

In [14]: xp = np.linspace(0,8,100).reshape(-1,1)
ypt = -0.8*xp + 3

plt.figure(figsize=(10, 6))
plt.plot(x1[C1], x2[C1], 'ro', label='C1')
plt.plot(x1[C2], x2[C2], 'bo', label='C2')
plt.plot(xp, ypt, '--k', label='True')
plt.title('linearly and strictly separable classes', fontweight = 'bold', fo
ntsize = 15)
plt.xlabel('$x_1$', fontsize = 20)
plt.ylabel('$x_2$', fontsize = 20)
plt.legend(loc = 4)
plt.xlim([0, 8])
plt.ylim([-4, 3])
plt.show()

```



- Given:
 - Hyperplane defined by ω and ω_0
 - Animals coordinates (or features) x

- Decision making:

$$\omega^T x + \omega_0 > 0 \implies x \text{ belongs to } C_1$$

$$\omega^T x + \omega_0 < 0 \implies x \text{ belongs to } C_2$$

- Find ω and ω_0 such that x given $\omega^T x + \omega_0 = 0$

or

- Find ω and ω_0 such that $x \in C_1$ given $\omega^T x + \omega_0 > 1$ and $x \in C_2$ given $\omega^T x + \omega_0 < -1$

$$\omega^T x + \omega_0 > b$$


$$\iff \frac{\omega^T}{b} x + \frac{\omega_0}{b} > 1$$

$$\iff \omega'^T x + \omega'_0 > 1$$

- Same problem if strictly separable

```
In [15]: # see how data are generated
xp = np.linspace(0,8,100).reshape(-1,1)
ypt = -0.8*xp + 3

plt.figure(figsize=(10, 6))
plt.plot(x1[C1], x2[C1], 'ro', label='C1')
plt.plot(x1[C2], x2[C2], 'bo', label='C2')
plt.plot(xp, ypt, '--k', label='true')
plt.plot(xp, ypt-1, '-k')
plt.plot(xp, ypt+1, '-k')
plt.title('linearly and strictly separable classes', fontweight = 'bold', fo
ntsize = 15)
plt.xlabel('$x_1$', fontsize = 20)
plt.ylabel('$x_2$', fontsize = 20)
plt.legend(loc = 4)
plt.xlim([0, 8])
plt.ylim([-4, 3])
plt.show()
```



3.2.1. Optimization Formulation 1

- $n (= 2)$ features
- $m = N + M$ data points in training set

$$x^{(i)} = \begin{bmatrix} x_1^{(i)} \\ x_2^{(i)} \end{bmatrix} \text{ with } \omega = \begin{bmatrix} \omega_1 \\ \omega_2 \end{bmatrix} \quad \text{or} \quad x^{(i)} = \begin{bmatrix} 1 \\ x_1^{(i)} \\ x_2^{(i)} \end{bmatrix} \text{ with } \omega = \begin{bmatrix} \omega_0 \\ \omega_1 \\ \omega_2 \end{bmatrix}$$

- N belongs to C_1 in training set
- M belongs to C_2 in training set
- ω and ω_0 are the unknown variables

minimize something

$$\text{subject to } \begin{cases} \omega^T x^{(1)} + \omega_0 \geq 1 \\ \omega^T x^{(2)} + \omega_0 \geq 1 \\ \vdots \\ \omega^T x^{(N)} + \omega_0 \geq 1 \end{cases} \quad \text{or}$$

$$\begin{cases} \omega^T x^{(N+1)} + \omega_0 \leq -1 \\ \omega^T x^{(N+2)} + \omega_0 \leq -1 \\ \vdots \\ \omega^T x^{(N+M)} + \omega_0 \leq -1 \end{cases}$$

minimize something

$$\text{subject to } \begin{cases} \omega^T x^{(1)} \geq 1 \\ \omega^T x^{(2)} \geq 1 \\ \vdots \\ \omega^T x^{(N)} \geq 1 \end{cases}$$

$$\begin{cases} \omega^T x^{(N+1)} \leq -1 \\ \omega^T x^{(N+2)} \leq -1 \\ \vdots \\ \omega^T x^{(N+M)} \leq -1 \end{cases}$$

Code (CVXPY)

$$X_1 = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(N)})^T \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} \\ \vdots & \vdots & \vdots \\ 1 & x_1^{(N)} & x_2^{(N)} \end{bmatrix}$$

$$X_2 = \begin{bmatrix} (x^{(N+1)})^T \\ (x^{(N+2)})^T \\ \vdots \\ (x^{(N+M)})^T \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(N+1)} & x_2^{(N+1)} \\ 1 & x_1^{(N+2)} & x_2^{(N+2)} \\ \vdots & \vdots & \vdots \\ 1 & x_1^{(N+M)} & x_2^{(N+M)} \end{bmatrix}$$

minimize something
subject to $X_1 \omega \geq 1$
 $X_2 \omega \leq -1$

minimize something
subject to $X_1 \omega \geq 1$
 $X_2 \omega \leq -1$

```
In [16]: # CVXPY using simple classification
import cvxpy as cvx

N = C1.shape[0]
M = C2.shape[0]

X1 = np.hstack([np.ones([N,1]), x1[C1], x2[C1]])
X2 = np.hstack([np.ones([M,1]), x1[C2], x2[C2]])


X1 = np.asmatrix(X1)
X2 = np.asmatrix(X2)
```

```
In [17]: w = cvx.Variable(3,1)
obj = cvx.Minimize(1)
const = [X1*w >= 1, X2*w <= -1]
prob = cvx.Problem(obj, const).solve()

w = w.value
```

```
In [18]: xp = np.linspace(0,8,100).reshape(-1,1)
yp = - w[1,0]/w[2,0]*xp - w[0,0]/w[2,0]

plt.figure(figsize=(10, 6))
plt.plot(X1[:,1], X1[:,2], 'ro', label='C1')
plt.plot(X2[:,1], X2[:,2], 'bo', label='C2')
plt.plot(xp, yp, 'k', label='SVM')
plt.plot(xp, ypt, '--k', label='true')
plt.xlim([0,8])
plt.xlabel('$x_1$', fontsize = 20)
plt.ylabel('$x_2$', fontsize = 20)
plt.legend(loc = 4, fontsize = 15)
plt.show()
```



3.2.2. Outlier

- Note that in the real world, you may have noise, errors, or outliers that do not accurately represent the actual phenomena
- Non-separable case
- No solutions (hyperplane) exist
 - We will allow some training examples to be misclassified !
 - but we want their number to be minimized

```


In [19]: X1 = np.hstack([np.ones([N,1]), x1[C1], x2[C1]])
X2 = np.hstack([np.ones([M,1]), x1[C2], x2[C2]])

outlier = np.array([1, 3, 2]).reshape(-1,1)
X2 = np.vstack([X2, outlier.T])

X1 = np.asmatrix(X1)
X2 = np.asmatrix(X2)

plt.figure(figsize=(10, 6))
plt.plot(X1[:,1], X1[:,2], 'ro', label='C1')
plt.plot(X2[:,1], X2[:,2], 'bo', label='C2')
plt.xlim([0,8])
plt.xlabel('$x_1$', fontsize = 20)
plt.ylabel('$x_2$', fontsize = 20)
plt.legend(loc = 4, fontsize = 15)
plt.show()

```



minimize something
 subject to $X_1\omega \geq 1$
 $X_2\omega \leq -1$

```

In [20]: w = cvx.Variable(3,1)
obj = cvx.Minimize(1)
const = [X1*w >= 1, X2*w <= -1]
prob = cvx.Problem(obj, const).solve()

print(w.value)

```

None

- No solutions (hyperplane) exist
- We will allow some training examples to be misclassified !
- but we want their number to be minimized

3.2.3. Optimization Formulation 2

- $n (= 2)$ features
- $m = N + M$ data points in a training set

$$x^i = \begin{bmatrix} 1 \\ x_1^{(i)} \\ x_2^{(i)} \end{bmatrix} \quad \text{with } \omega = \begin{bmatrix} \omega_0 \\ \omega_1 \\ \omega_2 \end{bmatrix} \quad \begin{array}{ll} \text{minimize} & \text{something} \\ \text{subject to} & X_1 \omega \geq 1 \\ & X_2 \omega \leq -1 \end{array}$$

- N belongs to C_1 in training set
- M belongs to C_2 in training set
- ω and ω_0 are the variables (unknown)
- For the non-separable case, we relax the above constraints
- Need slack variables u and v where all are positive

The optimization problem for the non-separable case

$$\begin{array}{ll} \text{minimize} & \sum_{i=1}^N u_i + \sum_{i=1}^M v_i \\ \text{subject to} & \begin{cases} \omega^T x^{(1)} \geq 1 - u_1 \\ \omega^T x^{(2)} \geq 1 - u_2 \\ \vdots \\ \omega^T x^{(N)} \geq 1 - u_N \end{cases} \\ & \begin{cases} \omega^T x^{(N+1)} \leq -(1 - v_1) \\ \omega^T x^{(N+2)} \leq -(1 - v_2) \\ \vdots \\ \omega^T x^{(N+M)} \leq -(1 - v_M) \end{cases} \\ & \begin{cases} u \geq 0 \\ v \geq 0 \end{cases} \end{array}$$

- Expressed in a matrix form

$$X_1 = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(N)})^T \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} \\ \vdots & \vdots & \vdots \\ 1 & x_1^{(N)} & x_2^{(N)} \end{bmatrix}$$

$$X_2 = \begin{bmatrix} (x^{(N+1)})^T \\ (x^{(N+2)})^T \\ \vdots \\ (x^{(N+M)})^T \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(N+1)} & x_2^{(N+1)} \\ 1 & x_1^{(N+2)} & x_2^{(N+2)} \\ \vdots & \vdots & \vdots \\ 1 & x_1^{(N+M)} & x_2^{(N+M)} \end{bmatrix}$$

$$u = \begin{bmatrix} u_1 \\ \vdots \\ u_N \end{bmatrix}$$

$$v = \begin{bmatrix} v_1 \\ \vdots \\ v_M \end{bmatrix}$$

$$\begin{array}{ll} \text{minimize} & 1^T u + 1^T v \\ \text{subject to} & X_1 \omega \geq 1 - u \\ & X_2 \omega \leq -(1 - v) \\ & u \geq 0 \\ & v \geq 0 \end{array}$$

```

In [21]: X1 = np.hstack([np.ones([C1.shape[0],1]), x1[C1], x2[C1]])
X2 = np.hstack([np.ones([C2.shape[0],1]), x1[C2], x2[C2]])

outlier = np.array([1, 2, 2]).reshape(-1,1)
X2 = np.vstack([X2, outlier.T])

X1 = np.asmatrix(X1)
X2 = np.asmatrix(X2)

N = X1.shape[0]
M = X2.shape[0]

w = cvx.Variable(3,1)
u = cvx.Variable(N,1)
v = cvx.Variable(M,1)
obj = cvx.Minimize(np.ones((1,N))*u + np.ones((1,M))*v)
const = [X1*w >= 1-u, X2*w <= -(1-v), u >= 0, v >= 0 ]
prob = cvx.Problem(obj, const).solve()

w = w.value


```

```

In [22]: xp = np.linspace(0,8,100).reshape(-1,1)
yp = - w[1,0]/w[2,0]*xp - w[0,0]/w[2,0]

plt.figure(figsize=(10, 6))
plt.plot(X1[:,1], X1[:,2], 'ro', label='C1')
plt.plot(X2[:,1], X2[:,2], 'bo', label='C2')
plt.plot(xp, yp, '--k', label='SVM')
plt.plot(xp, yp-1/w[2,0], '-k')
plt.plot(xp, yp+1/w[2,0], '-k')
plt.xlim([0,8])
plt.xlabel('$x_1$', fontsize = 20)
plt.ylabel('$x_2$', fontsize = 20)
plt.legend(loc = 4, fontsize = 15)
plt.show()


```



Further improvement

- Notice that hyperplane is not as accurately represent the division due to the outlier
- Can we do better when there are noise data or outliers?
- Yes, but we need to look beyond LP
- Idea: large margin leads to good generalization on the test data

3.3. Maximize Margin (Finally, it is Support Vector Machine)



- Distance (= margin)

$$\text{margin} = \frac{2}{\|\omega\|_2}$$

- Minimize $\|\omega\|_2$ to maximize the margin (closest samples from the decision line)

maximize {minimum distance}

- Use gamma (γ) as a weighting between the followings:
 - Bigger margin given robustness to outliers
 - Hyperplane that has few (or no) errors

$$\begin{aligned} &\text{minimize} && \|\omega\|_2 + \gamma(1^T u + 1^T v) \\ &\text{subject to} && X_1 \omega + \omega_0 \geq 1 - u \\ & && X_2 \omega + \omega_0 \leq -(1 - v) \\ & && u \geq 0 \\ & && v \geq 0 \end{aligned}$$

```

In [23]: g = 1
w = cvx.Variable(3,1)
u = cvx.Variable(N,1)
v = cvx.Variable(M,1)
obj = cvx.Minimize(cvx.norm(w,2) + g*(np.ones((1,N))*u + np.ones((1,M))*v))
const = [X1*w >= 1-u, X2*w <= -(1-v), u >= 0, v >= 0 ]
prob = cvx.Problem(obj, const).solve()

w = w.value


```

```

In [24]: xp = np.linspace(0,8,100).reshape(-1,1)
yp = - w[1,0]/w[2,0]*xp - w[0,0]/w[2,0]

plt.figure(figsize=(10, 6))
plt.plot(X1[:,1], X1[:,2], 'ro', label='C1')
plt.plot(X2[:,1], X2[:,2], 'bo', label='C2')
plt.plot(xp, yp, '--k', label='SVM')
plt.plot(xp, yp-1/w[2,0], '-k')
plt.plot(xp, yp+1/w[2,0], '-k')
plt.xlim([0,8])
plt.xlabel('$x_1$', fontsize = 20)
plt.ylabel('$x_2$', fontsize = 20)
plt.legend(loc = 4, fontsize = 15)
plt.show()

```




3.4. Logistic Regression

- Logistic regression is a classification algorithm
 - don't be confused
- Perceptron: make use of sign of data
- SVM: make use of margin (minimum distance)
 - Distance from a single data point
- We want to use distance information of ALL data points
 - logistic regression

Using Distances

- basic idea: to find the decision boundary (hyperplane) of $g(x) = \omega^T x = 0$ such that maximizes $\prod_i |h_i| \rightarrow$ optimization



- Inequality of arithmetic and geometric means

$$\frac{x_1 + x_2 + \dots + x_m}{m} \geq \sqrt[m]{x_1 \cdot x_2 \cdot \dots \cdot x_m}$$


and that equality holds if and only if $x_1 = x_2 = \dots = x_m$

- Roughly speaking, this optimization of $\max \prod_i |h_i|$ tends to position a hyperplane in the middle of two classes


$$h = \frac{g(x)}{\|\omega\|} = \frac{\omega^T x}{\|\omega\|} \sim \omega^T x$$

Using all Distances with Outliers

- SVM vs. Logistic Regression




SVM



Logistic Regression

Sigmoid function


- We link or squeeze $(-\infty, +\infty)$ to $(0, 1)$ for several reasons:



- If $\sigma(z)$ is the sigmoid function, or the logistic function

$$\sigma(z) = \frac{1}{1 + e^{-z}} \implies \sigma(\omega^T x) = \frac{1}{1 + e^{-\omega^T x}}$$

- logistic function always generates a value between 0 and 1
- Crosses 0.5 at the origin, then flattens out




- Benefit of mapping via the logistic function
 - monotonic: same or similar optimization solution
 - continuous and differentiable: good for gradient descent optimization
 - probability or confidence: can be considered as probability

$$P(y = +1 \mid x, \omega) = \frac{1}{1 + e^{-\omega^T x}} \in [0, 1]$$

- Goal: we need to fit ω to our data

$$\max \prod_i |h_i|$$

- Classified based on probability



```
In [25]: m = 200

X0 = np.random.multivariate_normal([0, 0], np.eye(2), m)
X1 = np.random.multivariate_normal([10, 10], np.eye(2), m)

X = np.vstack([X0, X1])
y = np.vstack([np.zeros([m,1]), np.ones([m,1])])
```

```
In [26]: from sklearn import linear_model

clf = linear_model.LogisticRegression()
clf.fit(X, np.ravel(y))
```

```
Out[26]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)
```

```
In [27]: X_new = np.array([2, 0]).reshape(1, -1)
pred = clf.predict(X_new)

print(pred)

[0.]
```


```
In [28]: pred = clf.predict_proba(X_new)

print(pred)

[[0.9407617 0.0592383]]
```

```
In [29]: plt.figure(figsize=(10, 6))
plt.plot(X0[:,0], X0[:,1], '.b', label='Class 0')
plt.plot(X1[:,0], X1[:,1], '.k', label='Class 1')
plt.plot(X_new[0,0], X_new[0,1], 'o', label='New Data', ms=5, mew=5)

plt.title('Logistic Regression', fontsize=15)
plt.legend(loc='lower right', fontsize=15)
plt.xlabel('X1', fontsize=15)
plt.ylabel('X2', fontsize=15)
plt.grid(alpha=0.3)
plt.axis('equal')
plt.show()
```



```
In [30]: %%javascript
$.getScript('https://kmahelona.github.io/ipython_notebook_goodies/ipython_notebook_toc.js')
```