

# Supervised Learning

with Scikit Learn

by Prof. Seungchul Lee  
Industrial AI Lab  
<http://isystems.unist.ac.kr/>  
POSTECH


## Table of Contents

- I. 1. Supervised learning
- II. 2. Regression
  - I. 2.1. k-Nearest Neighbor Regression
  - II. 1.2. Linear Regression
- III. 2. Classification
  - I. 2.1. Data Generation for Classification
  - II. 2.2. K-nearest neighbors
- IV. 3. Support Vector Machine (SVM)
- V. 4. Logistic Regression
- VI. 5. Nonlinear Classification

# 1. Supervised learning

- Given training set  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$
- Want to find a function  $g_\omega$  with learning parameter,  $\omega$ 
  - $g_\omega$  desired to be as close as possible to  $y$  for future  $(x, y)$
  - *i. e.* ,  $g_\omega(x) \sim y$
- Define a loss function  $\ell$
- Solve the following optimization problem:

$$\begin{aligned} \text{minimize} \quad & f(\omega) = \frac{1}{m} \sum_{i=1}^m \ell(g_\omega(x^{(i)}), y^{(i)}) \\ \text{subject to} \quad & \omega \in \omega \end{aligned}$$



## 2. Regression

### 2.1. k-Nearest Neighbor Regression

The goal is to make quantitative (real valued) predictions on the basis of a (vector of) features or attributes.

We write our model as


$$Y = f(X) + \epsilon$$

where  $\epsilon$  captures measurement errors and other discrepancies.

Then, with a good  $f$  we can make predictions of  $Y$  at new points  $X = x$ . One possible way so called "nearest neighbor method" is:

$$\hat{f} = \text{Ave} (Y \mid X \in \mathcal{N}(x))$$

where  $\mathcal{N}(x)$  is some neighborhood of  $x$




```
In [1]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

N = 100
w1 = 3
w0 = 2
x = np.random.uniform(0, 10, (N,1))
y = w1*x + w0 + 5*np.random.normal(0, 1, (N,1))
```

```
In [2]: plt.figure(figsize=(10, 6))
plt.plot(x, y, '.')

plt.title('Data Set', fontsize=15)
plt.xlabel('X', fontsize=15)
plt.ylabel('Y', fontsize=15)
plt.xlim([0, 10])
plt.grid(alpha=0.3)
plt.show()
```



```
In [3]: from sklearn.neighbors import KNeighborsRegressor
```

```
reg = KNeighborsRegressor(n_neighbors=10)
reg.fit(x, y)
```

```
Out[3]: KNeighborsRegressor(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=1, n_neighbors=10, p=2,
                             weights='uniform')
```

```
In [4]: pred = reg.predict(5)
print(pred)
```

```
[[15.98874782]]
```


```

In [5]: xp = np.linspace(0, 10, 100).reshape(-1, 1)
        yp = reg.predict(xp)

        plt.figure(figsize=(10, 6))
        plt.plot(x, y, '.', label='Original Data')
        plt.plot(xp, yp, 'r', label='Regression Result')

        plt.title('k-Nearest Neighbor Regression', fontsize=15)
        plt.xlabel('X', fontsize=15)
        plt.ylabel('Y', fontsize=15)
        plt.legend(loc=2, fontsize=15)
        plt.xlim([0, 10])
        plt.ylim([-5, 45])
        plt.grid(alpha=0.3)
        plt.show()

```



## 1.2. Linear Regression

Given  $\begin{cases} x_i : \text{inputs} \\ y_i : \text{outputs} \end{cases}$ , find  $\theta_1$  and  $\theta_2$

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix}, \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix} \approx \hat{y}_i = \theta_1 x_i + \theta_2$$


- $\hat{y}_i$  : predicted output

- $\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$  : Model parameters

$$\hat{y}_i = f(x_i, \theta) \text{ in general}$$

- in many cases, a linear model to predict  $y_i$  can be used

$$\hat{y}_i = \theta_1 x_i + \theta_2 \quad \text{such that} \quad \min_{\theta_1, \theta_2} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$




```

In [6]: N = 100
        w1 = 3
        w0 = 2
        x = np.random.uniform(0, 10, (N,1))
        y = w1*x + w0 + 3*np.random.normal(0, 1, (N,1))

        plt.figure(figsize=(10, 6))
        plt.plot(x, y, 'b.')

        plt.title('Data Set', fontsize=15)
        plt.xlabel('X', fontsize=15)
        plt.ylabel('Y', fontsize=15)
        plt.xlim([0, 10])
        plt.grid(alpha=0.3)
        plt.show()

```



```

In [7]: from sklearn.linear_model import LinearRegression

```

```

        reg = LinearRegression()
        reg.fit(x, y)

```

```

Out[7]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

```

```

In [8]: pred = reg.predict(5)
        print(pred)

```

```


[[17.17945416]]

```

```
In [9]: xp = np.linspace(0, 10).reshape(-1,1)
yp = reg.predict(xp)

plt.figure(figsize=(10, 6))
plt.plot(x, y, '.', label='Original Data')
plt.plot(xp, yp, 'r', label='Regression Result')

plt.title('Linear Regression', fontsize=15)
plt.xlabel('X', fontsize=15)
plt.ylabel('Y', fontsize=15)
plt.legend(fontsize=15)
plt.xlim([0, 10])
plt.ylim([-5, 45])
plt.grid(alpha=0.3)
plt.show()
```



## 2. Classification

### 2.1. Data Generation for Classification

```
In [10]: import matplotlib.pyplot as plt

m = 200


X0 = np.random.multivariate_normal([0, 0], np.eye(2), m)
X1 = np.random.multivariate_normal([10, 10], np.eye(2), m)

X = np.vstack([X0, X1])
y = np.vstack([np.zeros([m,1]), np.ones([m,1])])
```



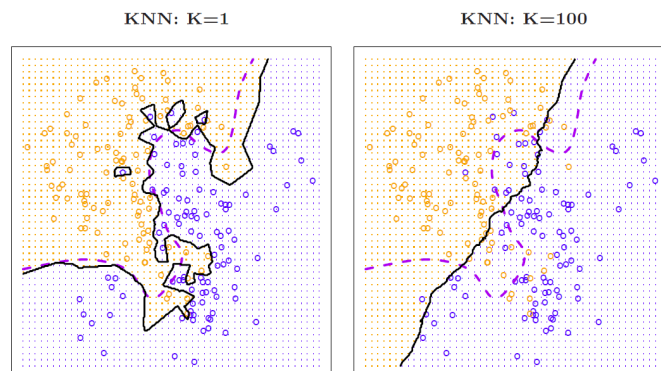
```
In [11]: plt.figure(figsize=(10, 6))
plt.plot(X0[:,0], X0[:,1], '.b', label='Class 0')
plt.plot(X1[:,0], X1[:,1], '.k', label='Class 1')

plt.title('Data Classes', fontsize=15)
plt.legend(loc='lower right', fontsize=15)
plt.xlabel('X1', fontsize=15)
plt.ylabel('X2', fontsize=15)
plt.axis('equal')
plt.grid(alpha=0.3)
plt.show()
```




## 2.2. K-nearest neighbors

- In k-NN classification, an object is assigned to the class most common among its  $k$  nearest neighbors ( $k$  is a positive integer, typically small).
- If  $k = 1$ , then the object is simply assigned to the class of that single nearest neighbor.



- Zoom in,



```
In [12]: from sklearn.neighbors import KNeighborsClassifier
```

```
clf = KNeighborsClassifier(n_neighbors=2)
clf.fit(X, np.ravel(y))
```

```
Out[12]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=1, n_neighbors=2, p=2,
                             weights='uniform')
```

```
In [13]: X_new = np.array([2,0]).reshape(1,-1)
pred = clf.predict(X_new)
```

```
print(pred)
```

```
[0.]
```


### 3. Support Vector Machine (SVM)

- 가장 많이 쓰이는 모델
- 경계선과 데이터 사이의 거리 (margin) 을 최대화 하는 모델
- Distance (= margin)

$$\text{margin} = \frac{2}{\|\omega\|_2}$$

- Minimize  $\|\omega\|_2$  to maximize the margin

$$\begin{aligned} &\text{minimize} && \|\omega\|_2 + \gamma(1^T u + 1^T v) \\ &\text{subject to} && X_1 \omega + \omega_0 \geq 1 - u \\ & && X_2 \omega + \omega_0 \leq -(1 - v) \\ & && u \geq 0 \\ & && v \geq 0 \end{aligned}$$



```
In [14]: from sklearn.svm import SVC
```

```
clf = SVC()
clf.fit(X, np.ravel(y))
```

```
Out[14]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```


```
In [15]: X_new = np.array([7, 10]).reshape(1, -1)
pred = clf.predict(X_new)
```

```
print(pred)
```

```
[1.]
```

## 4. Logistic Regression

- basic idea: to find the decision boundary (hyperplane) of  $g(x) = \omega^T x = 0$  such that maximizes  $\prod_i |h_i| \rightarrow$  optimization




- Inequality of arithmetic and geometric means

$$\frac{x_1 + x_2 + \dots + x_m}{m} \geq \sqrt[m]{x_1 \cdot x_2 \cdot \dots \cdot x_m}$$

and that equality holds if and only if  $x_1 = x_2 = \dots = x_m$

- Classified based on probability



```
In [16]: m = 200

X0 = np.random.multivariate_normal([0, 0], np.eye(2), m)
X1 = np.random.multivariate_normal([10, 10], np.eye(2), m)

X = np.vstack([X0, X1])
y = np.vstack([np.zeros([m,1]), np.ones([m,1])])
```

```
In [17]: from sklearn import linear_model
```

```
clf = linear_model.LogisticRegression()  
clf.fit(X, np.ravel(y))
```

```
Out[17]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,  
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,  
                             verbose=0, warm_start=False)
```

```
In [18]: X_new = np.array([2, 0]).reshape(1, -1)  
pred = clf.predict(X_new)
```

```
print(pred)
```

```
[0.]
```


```
In [19]: pred = clf.predict_proba(X_new)
```

```
print(pred)
```

```
[[0.94680429 0.05319571]]
```

```
In [20]: plt.figure(figsize=(10, 6))  
plt.plot(X0[:,0], X0[:,1], '.b', label='Class 0')  
plt.plot(X1[:,0], X1[:,1], '.k', label='Class 1')  
plt.plot(X_new[0,0], X_new[0,1], 'o', label='New Data', ms=5, mew=5)
```

```
plt.title('Logistic Regression', fontsize=15)  
plt.legend(loc='lower right', fontsize=15)  
plt.xlabel('X1', fontsize=15)  
plt.ylabel('X2', fontsize=15)  
plt.grid(alpha=0.3)  
plt.axis('equal')  
plt.show()
```




# 5. Nonlinear Classification

## Classifying non-linear separable data


- Consider the binary classification problem
  - each example represented by a single feature  $x$
  - No linear separator exists for this data




- Now map each example as  $x \rightarrow \{x, x^2\}$
- Data now becomes linearly separable in the new representation



- Linear in the new representation = nonlinear in the old representation
- Let's look at another example
  - Each example defined by a two features  $x = \{x_1, x_2\}$
  - No linear separator exists for this data



- Now map each example as  $x = \{x_1, x_2\} \rightarrow z = \{x_1^2, \sqrt{2}x_1x_2, x_2^2\}$ 
  - Each example now has three features (derived from the old representation)
- Data now becomes linear separable in the new representation



In [21]: `%%html  
<center><iframe src="https://www.youtube.com/embed/3liCbRZPrZA?rel=0"  
width="420" height="315" frameborder="0" allowfullscreen></iframe></center>`

### SVM with polynomial kernel visualization




- 이 부분 코드는 이해할 필요가 없으며, 개념적인 것만 이해하시면 됩니다
- Nonlinear Example

```

In [22]: N = 250 # number of points per class
D = 2 # dimensionality
K = 3 # number of classes
X = np.zeros([N*K, D]) # data matrix (each row = single example)
y = np.zeros(N*K) # class labels
for j in range(K):
    ix = range(N*j,N*(j+1))
    r = np.linspace(0.0, 1, N) # radius
    t = np.linspace(j*4, (j+1)*4, N) + np.random.randn(N)*0.2 # theta
    X[ix] = np.c_[r*np.sin(t), r*np.cos(t)]
    y[ix] = j

plt.figure(figsize=(6, 6))
plt.title('Nonlinear Data', fontsize=15)
plt.plot(X[y==0,0], X[y==0,1], '.', label='Class 0')
plt.plot(X[y==1,0], X[y==1,1], '.', label='Class 1')
plt.plot(X[y==2,0], X[y==2,1], '.', label='Class 2')
plt.xlim(min(X[:,0]) - 0.1, max(X[:,0]) + 0.1)
plt.ylim(min(X[:,1]) - 0.1, max(X[:,1]) + 0.1)
plt.legend(loc='lower right', fontsize=15)
plt.xlabel('X1', fontsize=15)
plt.ylabel('X2', fontsize=15)
plt.grid(alpha=0.3)
plt.show()

```



```

In [23]: from sklearn.svm import SVC

svc = SVC(kernel='linear', C=1).fit(X, y)
rbf_svc = SVC(kernel='rbf', C=1, gamma=5).fit(X, y)

```



```
In [24]: # create a mesh to plot in
```

```
h = 0.02 # step size in the mesh
x_min, x_max = X[:,0].min() - 0.1, X[:,0].max() + 0.1
y_min, y_max = X[:,1].min() - 0.1, X[:,1].max() + 0.1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))
```

```
In [25]: # title for the plots
```

```
titles = ['Linear Model', 'Nonlinear Model']
```

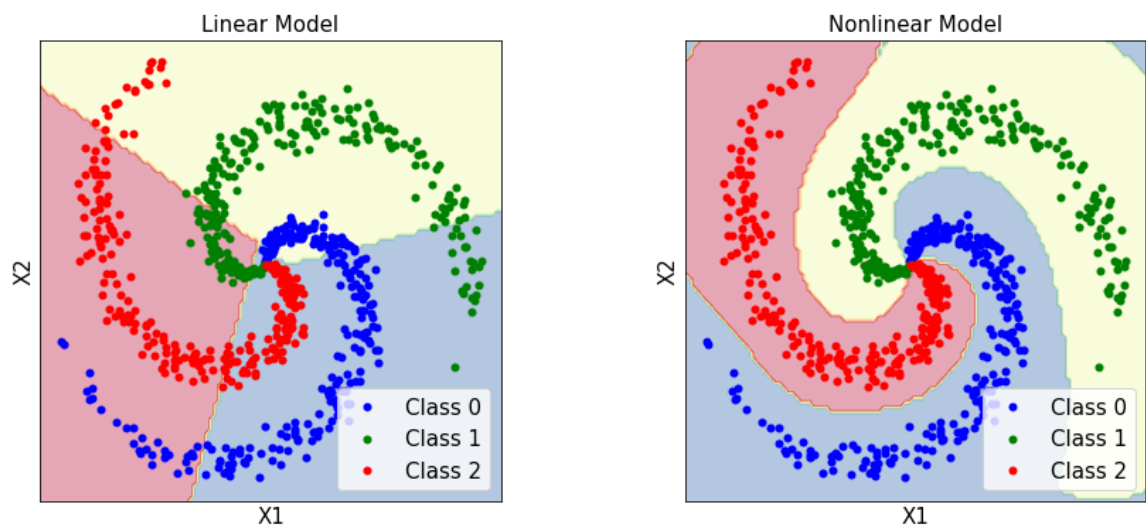
```
fig = plt.figure(figsize=(14, 6))
for i, clf in enumerate((svc, rbf_svc)):
    plt.subplot(1, 2, i+1)
    plt.subplots_adjust(wspace=0.4, hspace=0.4)

    Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])

    # Put the result into a color plot
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, cmap=plt.cm.Spectral_r, alpha=0.4)

    # Plot also the training points
    plt.plot(X[y==0,0], X[y==0,1], 'b.', label='Class 0', mew=3)
    plt.plot(X[y==1,0], X[y==1,1], 'g.', label='Class 1', mew=3)
    plt.plot(X[y==2,0], X[y==2,1], 'r.', label='Class 2', mew=3)
    plt.legend(loc='lower right', fontsize=15)
    plt.xlabel('X1', fontsize=15)
    plt.ylabel('X2', fontsize=15)
    plt.xlim(xx.min(), xx.max())
    plt.ylim(yy.min(), yy.max())
    plt.xticks([])
    plt.yticks([])
    plt.title(titles[i], fontsize=15)

plt.show()
```



```
In [26]: %%javascript
```

```
$.getScript('https://kmahelona.github.io/ipython_notebook_goodies/ipython_no  
tebook_toc.js')
```